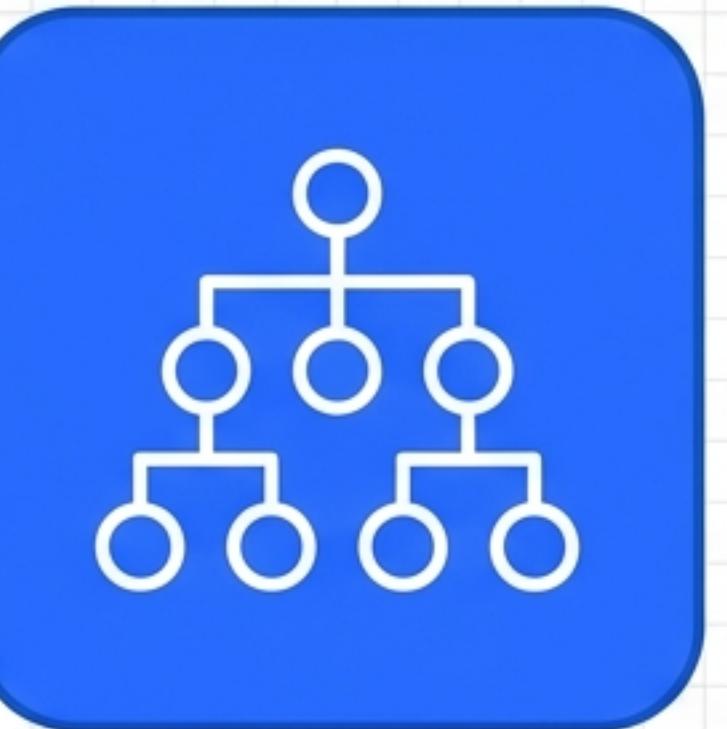


THE SYSTEM DESIGN ATLAS

A Field Guide to Modern Backend Architecture



INTENT:

Knowledge Transfer
& Reference

AUDIENCE:

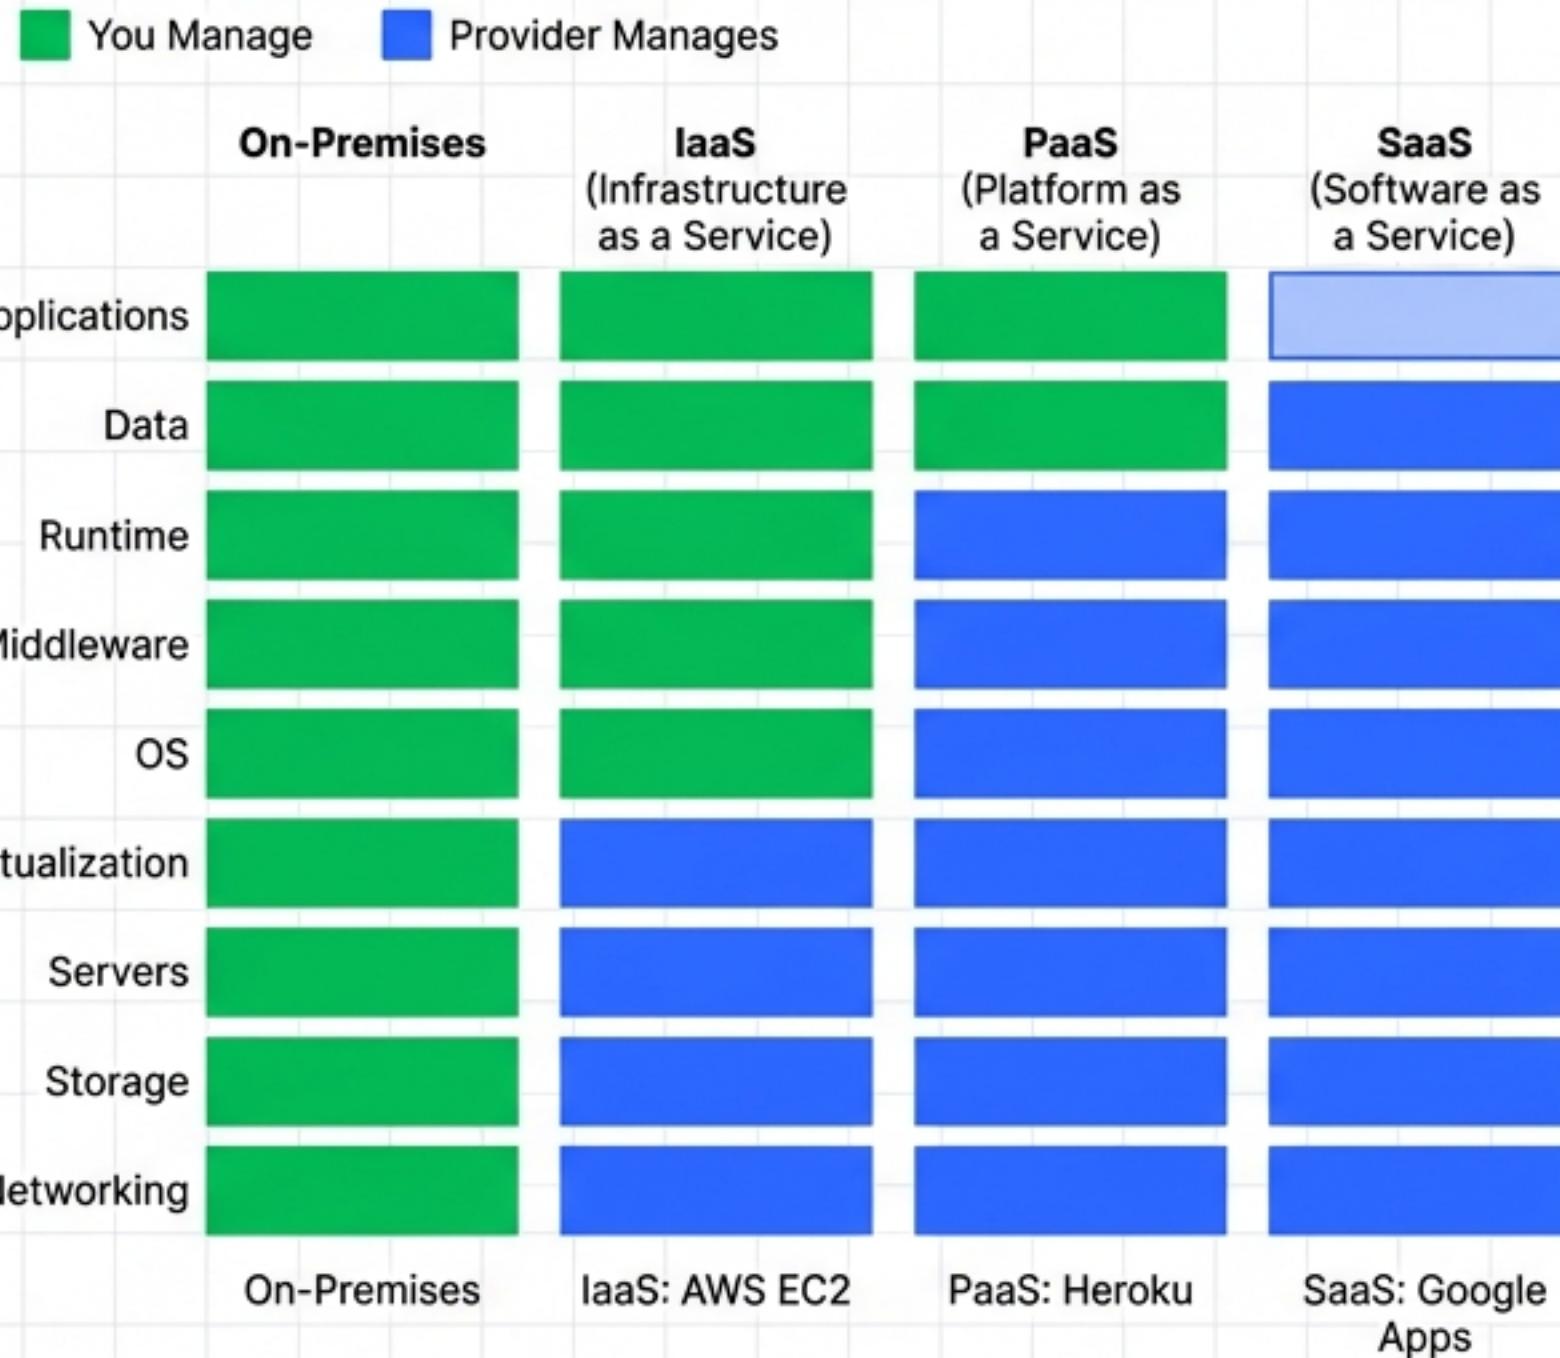
Engineers, Architects,
Managers

SCOPE:

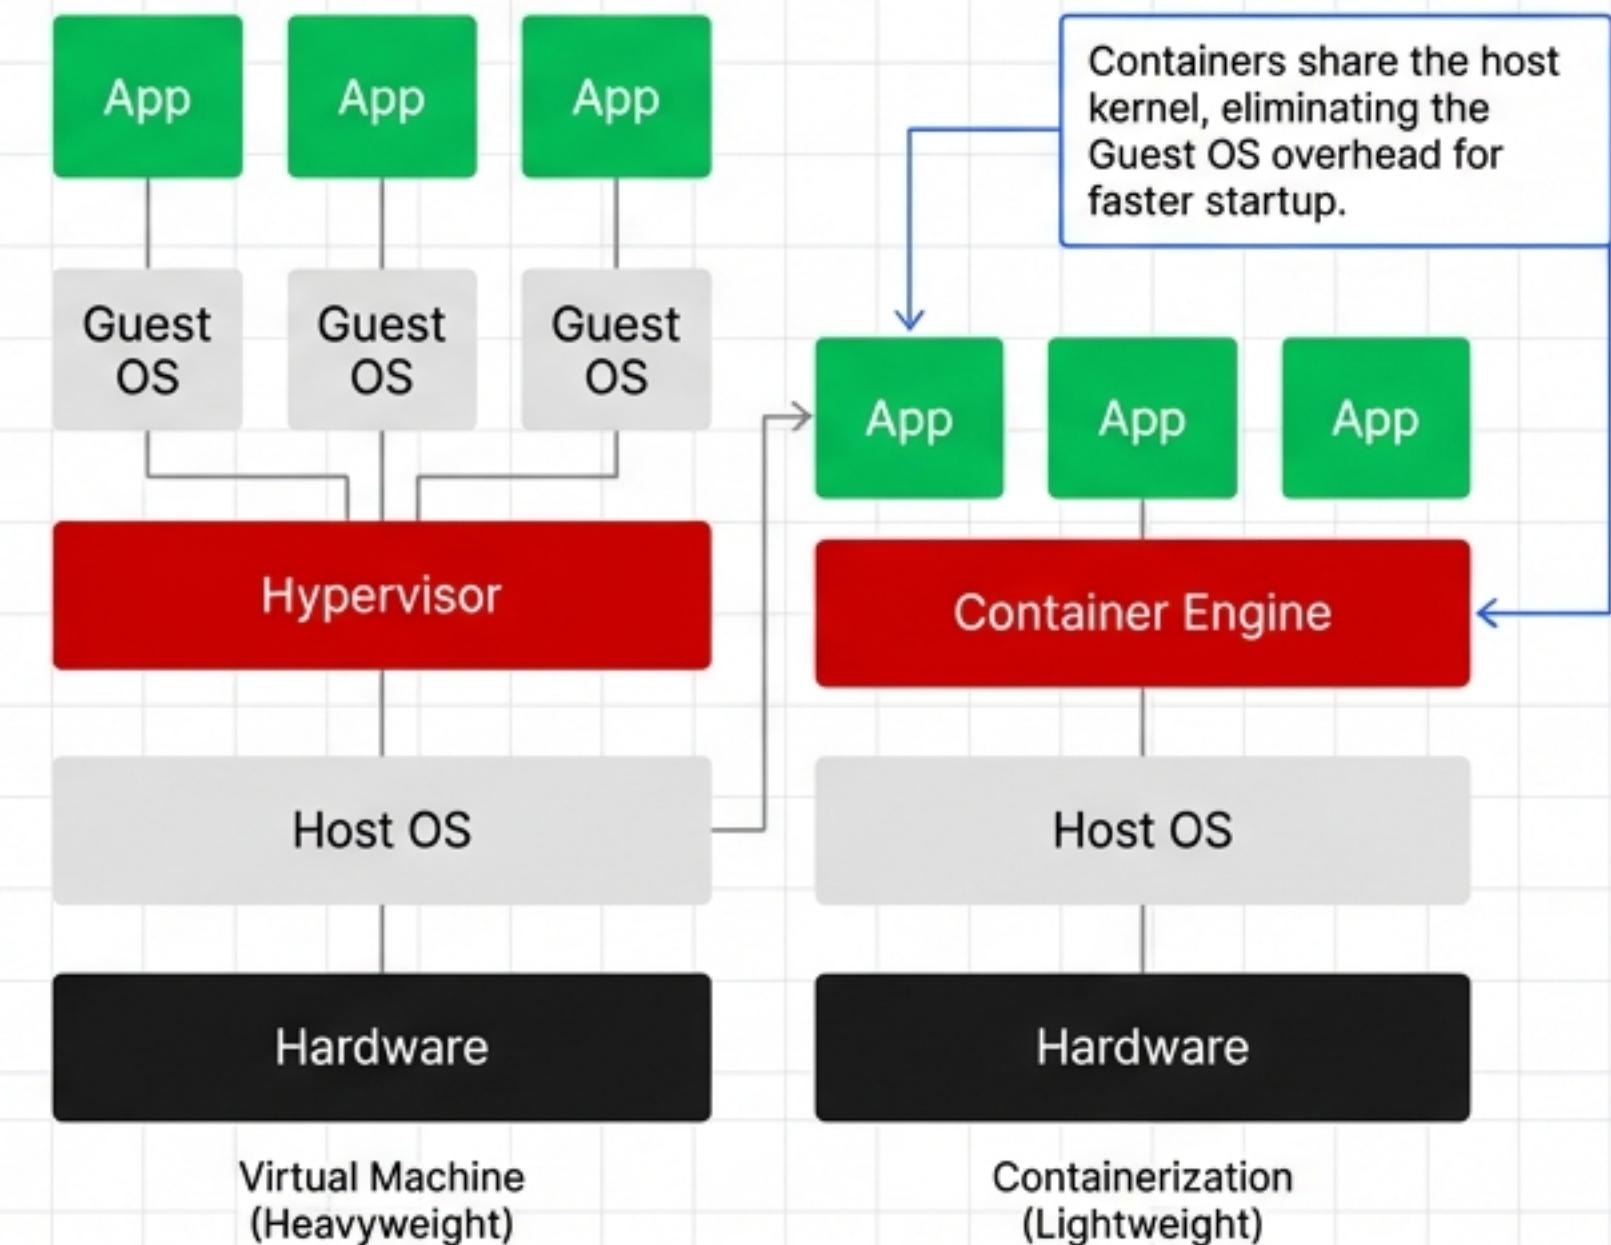
Infrastructure · Data · Network ·
Patterns · Case Studies

01. THE LANDSCAPE: CLOUD & EXECUTION

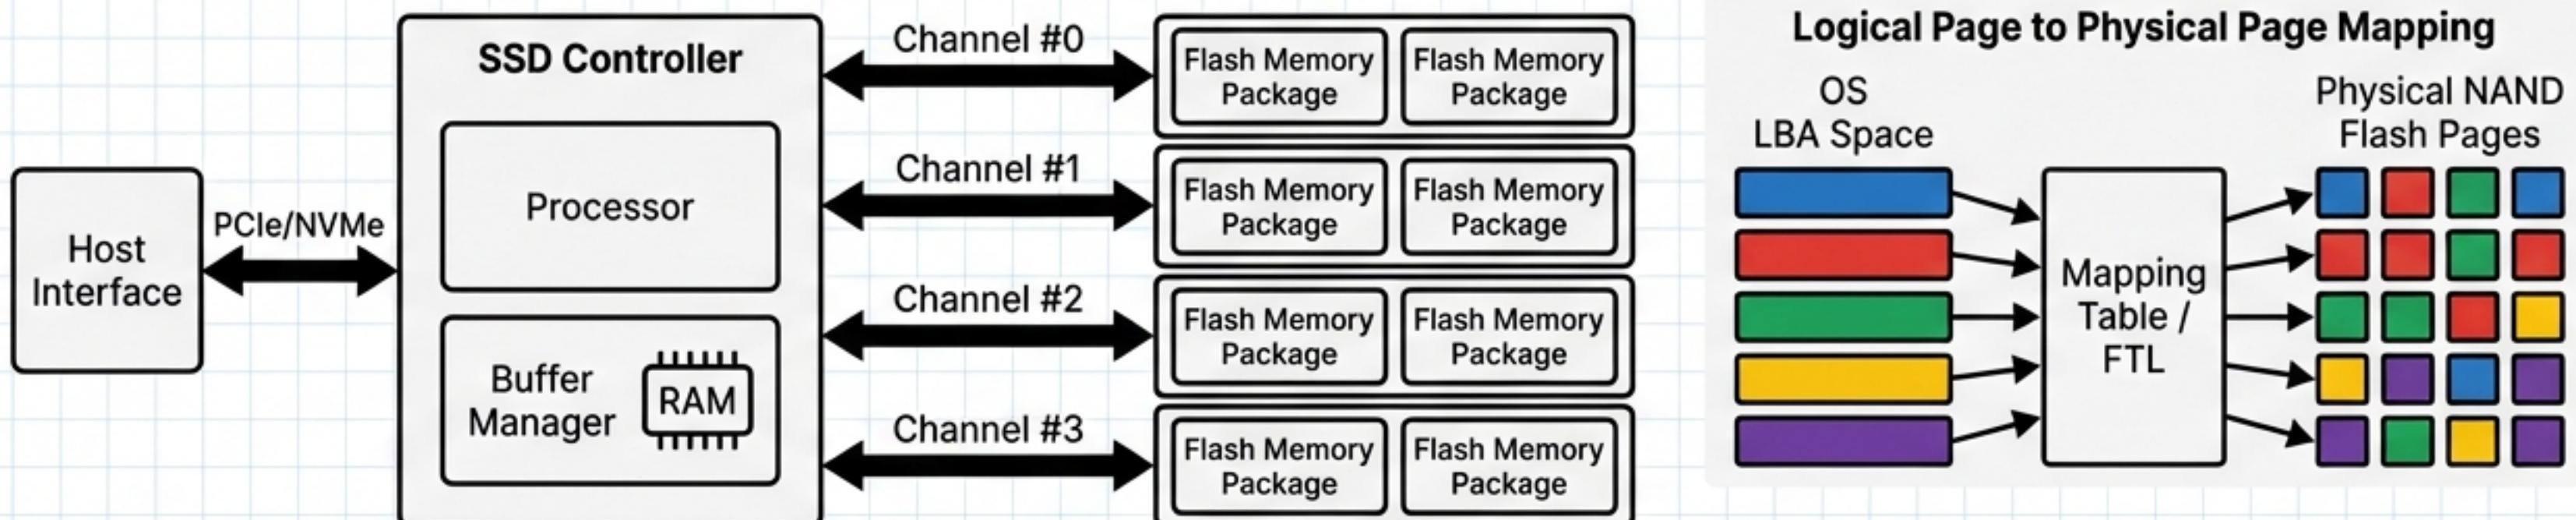
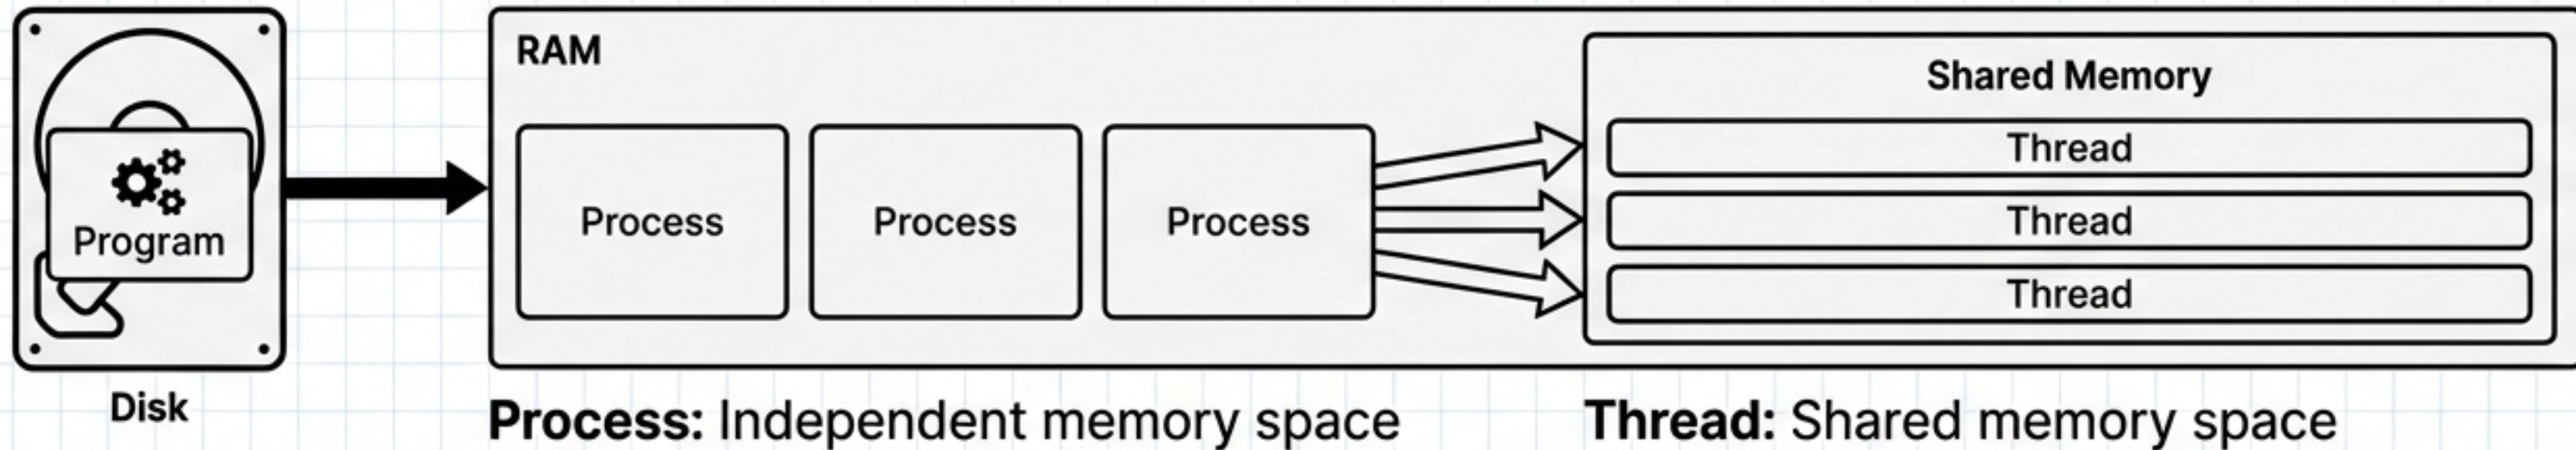
THE CLOUD RESPONSIBILITY MATRIX



VIRTUALIZATION VS. CONTAINERIZATION

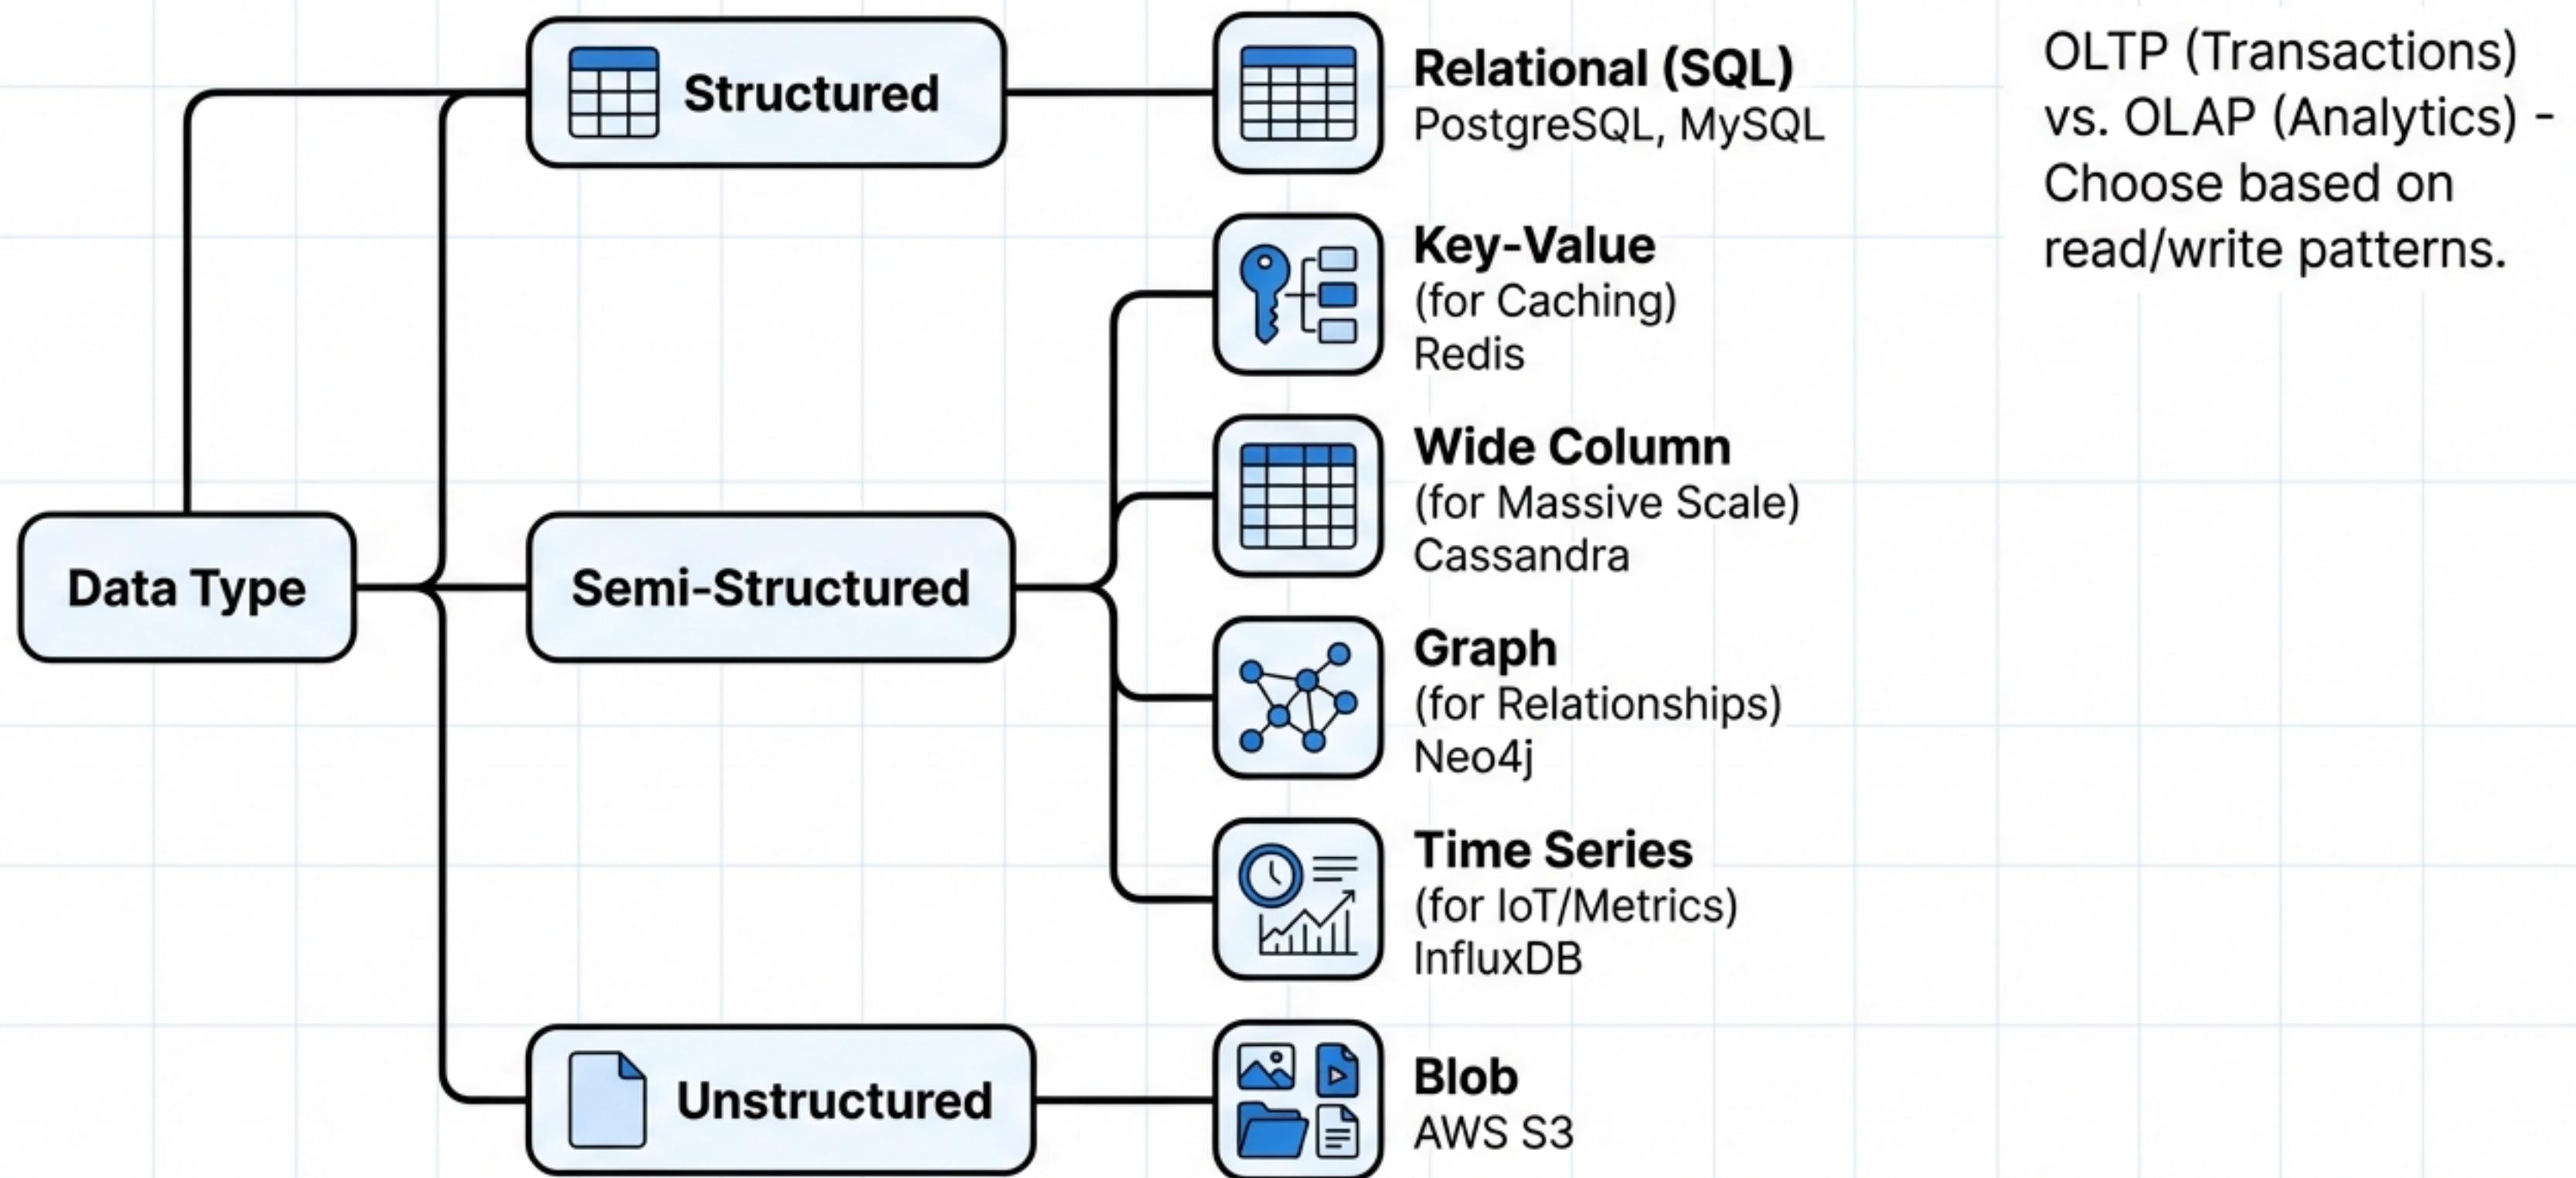


02. UNDER THE HOOD: THREADS & SSDs



Parallelism: The controller writes to multiple flash particles simultaneously via independent channels.

03. STRATEGY: CHOOSING A DATABASE



04. CONSISTENCY & ISOLATION

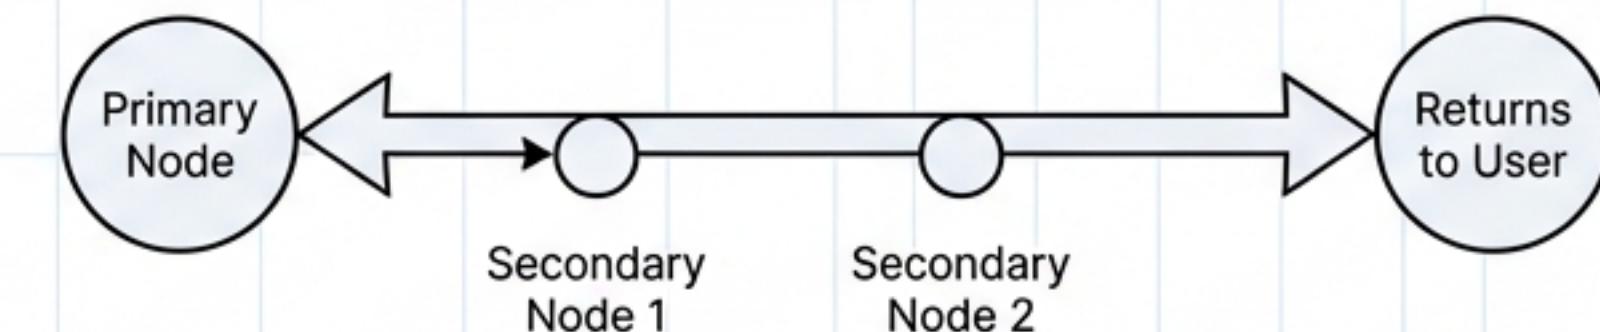
The Isolation Matrix

Isolation Level	Dirty Read	Non-repeatable Read	Phantom Read
Serializable (Strict)	Impossible	Impossible	Impossible
Repeatable Read	Impossible	Impossible	Probably
Read Committed	Impossible	Probably	Probably
Read Uncommitted	Probably	Probably	Probably

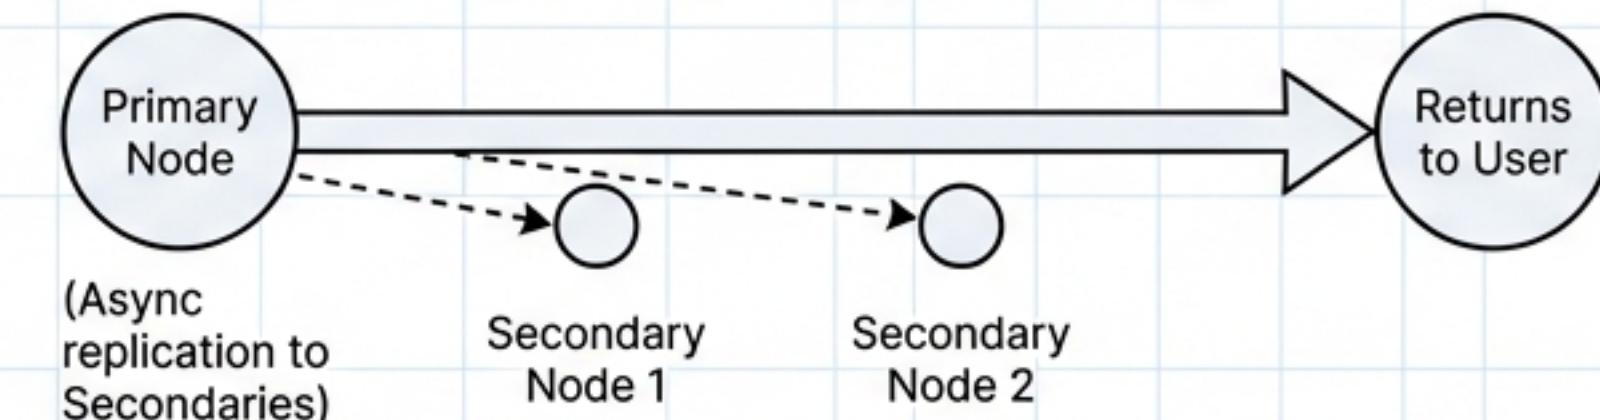
MVCC (Multi-Version Consistency Control) uses Transaction IDs to create snapshot views.

Consistency vs. Latency Trade-off

Scenario 1: High Consistency / High Latency



Scenario 2: Eventual Consistency / Low Latency

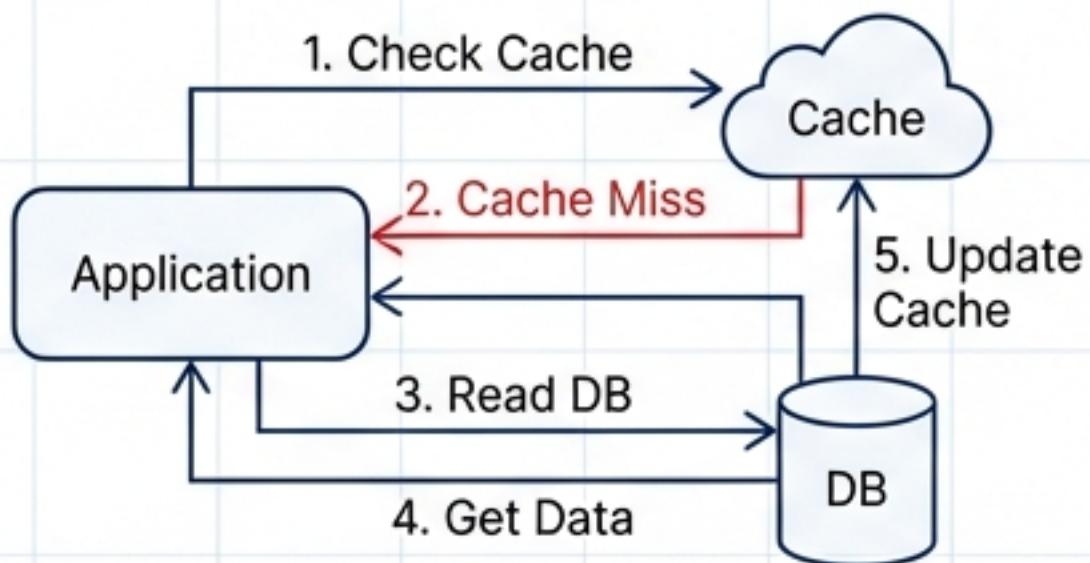


The CAP Theorem dictates you cannot have perfect Consistency and Availability simultaneously in a partitioned system.

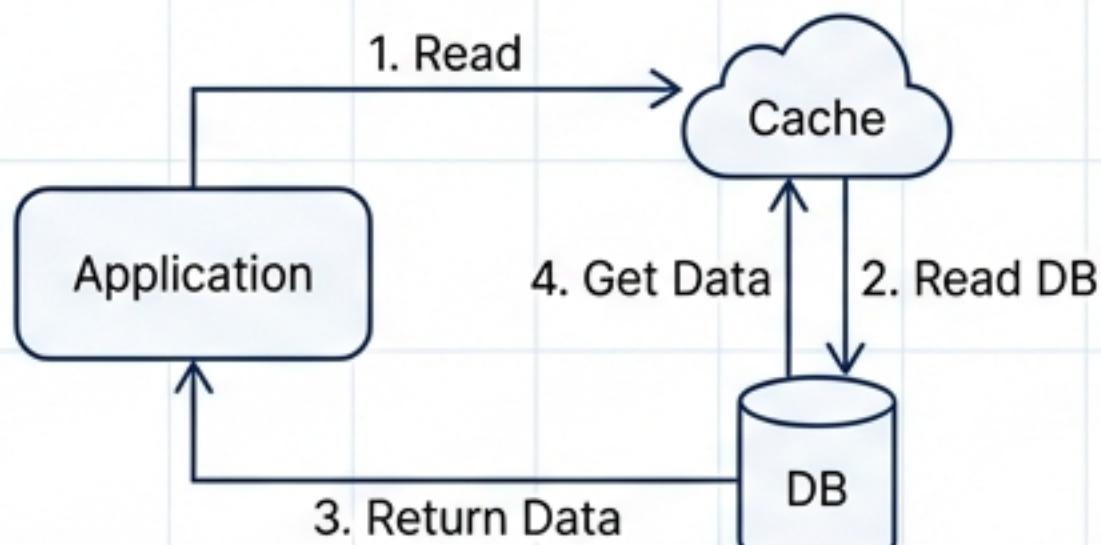
05. CACHING STRATEGIES

Caching Patterns

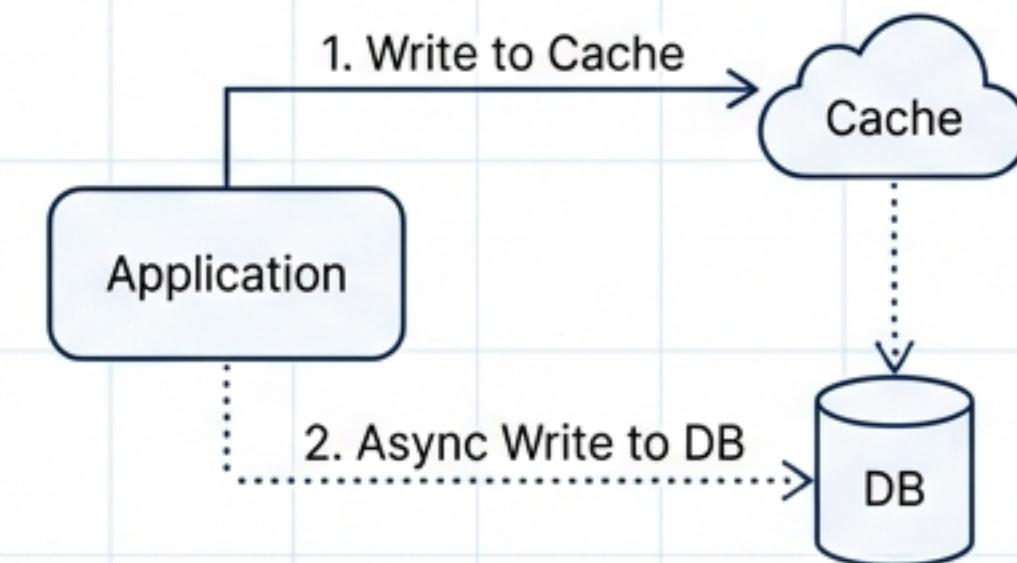
Cache-Aside



Read-Through



Write-Back

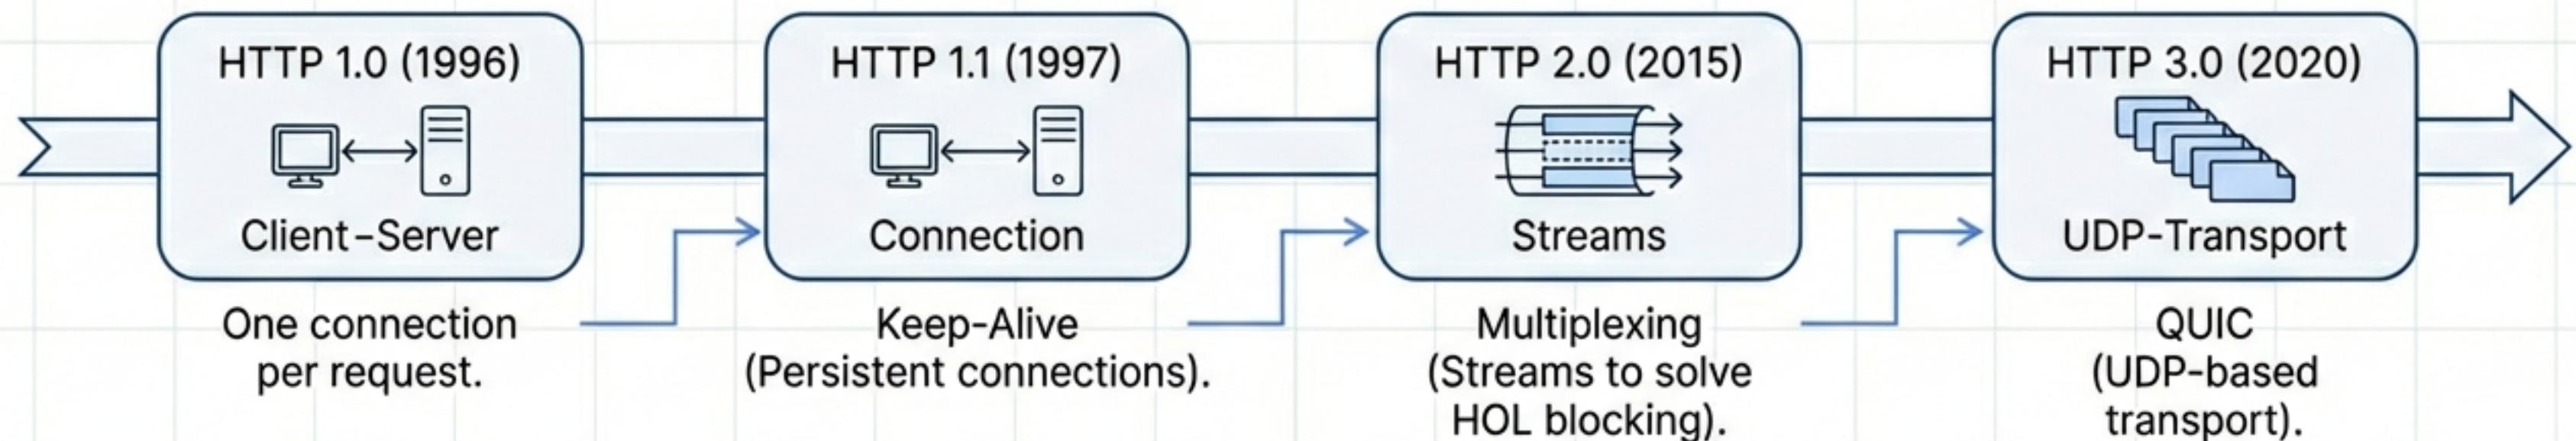


Redis vs. Memcached

Feature	Redis	Memcached
Data Structures	Complex (Lists, Sets, Sorted Sets, Bitmaps)	Simple Strings
Architecture	Single-Threaded Event Loop	Multi-Threaded
Persistence	Yes (RDB/AOF Snapshots)	No (In-memory only)
Eviction	Advanced Policies (LFU/LRU)	LRU Only

06. NETWORK: HTTP EVOLUTION & API STYLES

HTTP Evolution Timeline

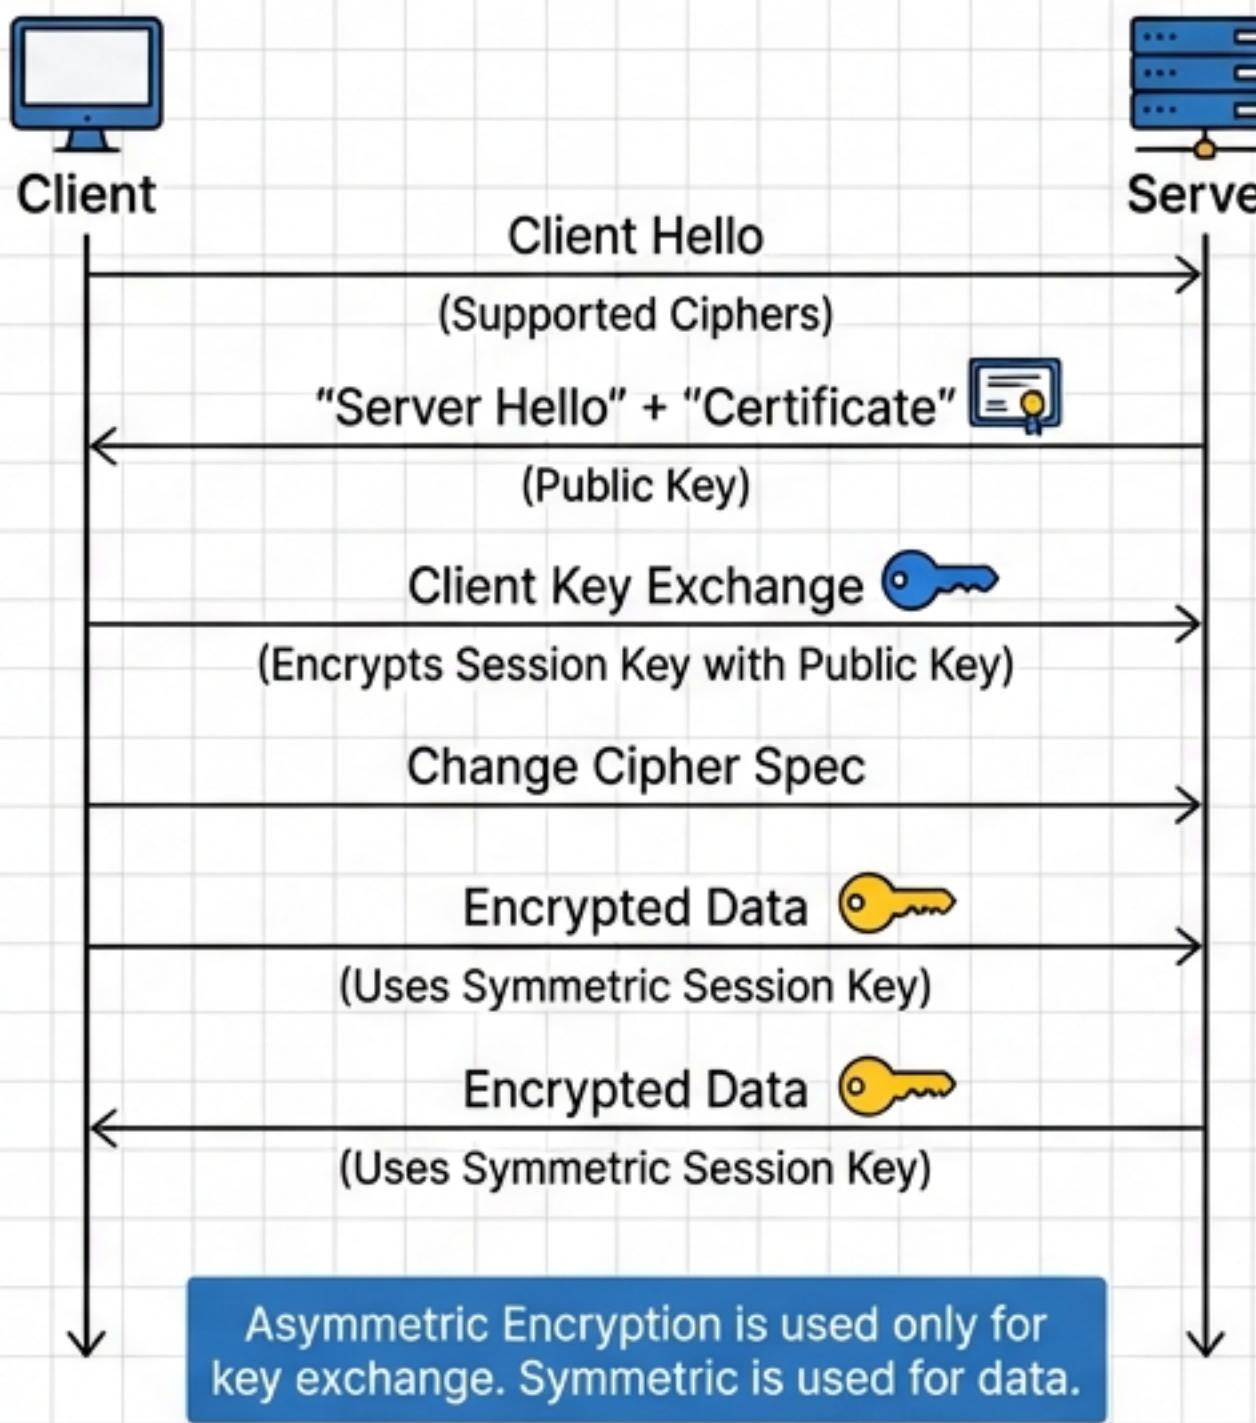


API Paradigm Matrix

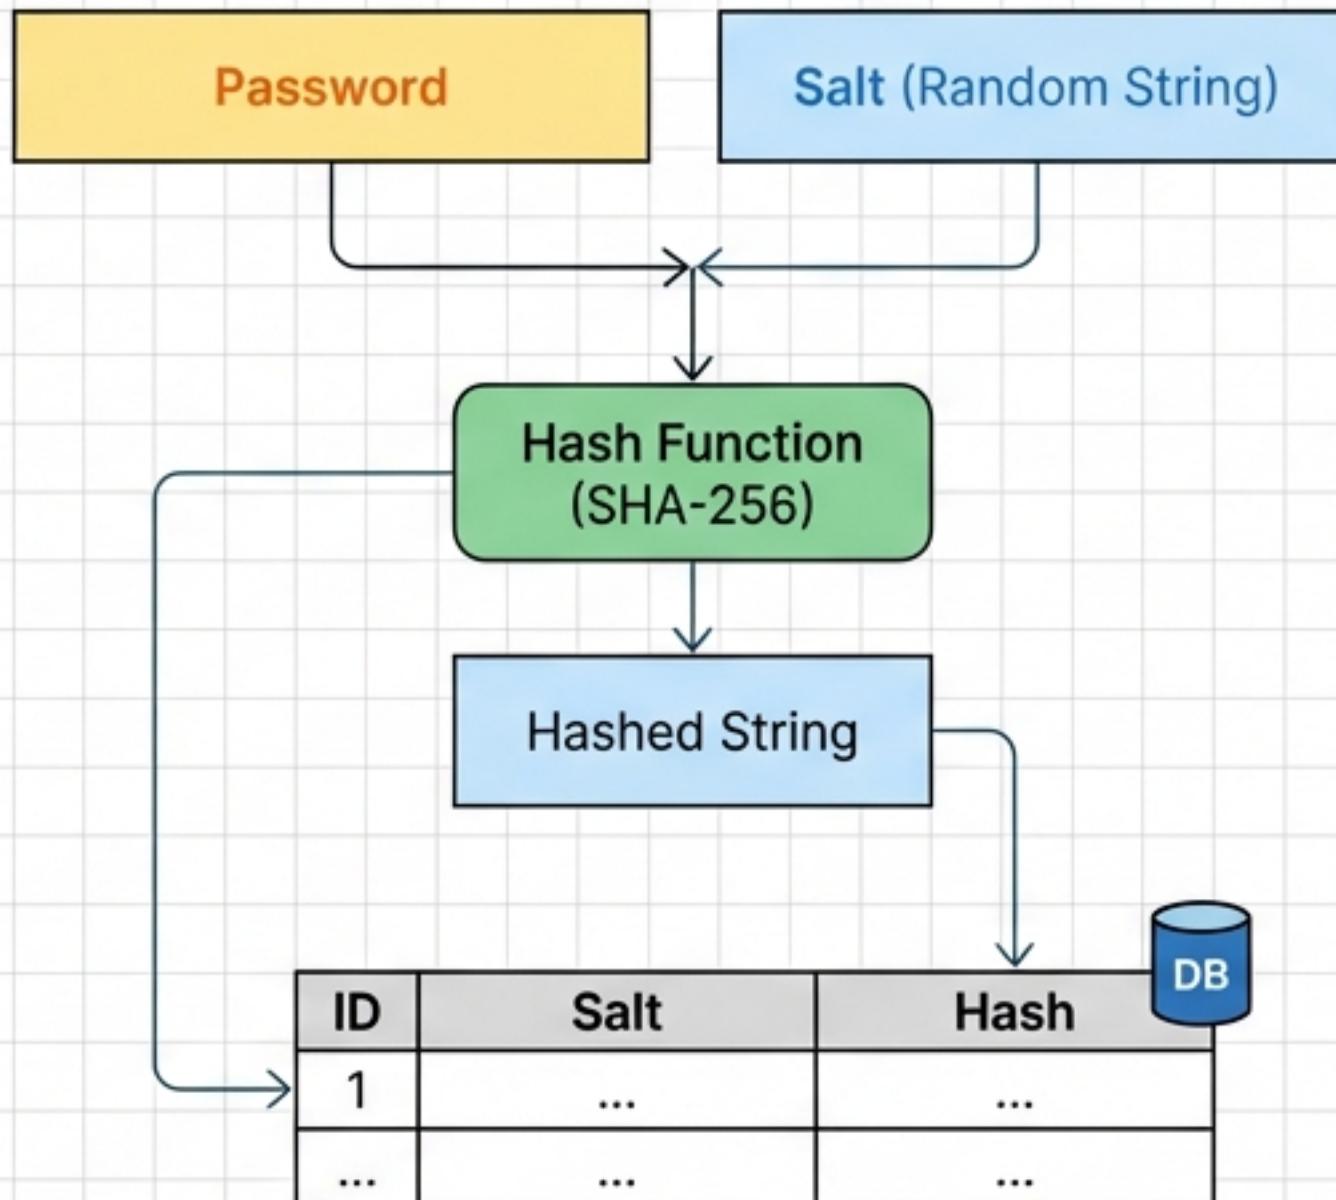
SOAP	REST	GraphQL	RPC (gRPC)
XML-only Strict Standard	Resource-based JSON/Stateless	Schema-driven Client-defined queries	Action-oriented Protobufs
Use case: Legacy Enterprise/Banking	Use case: Public Web APIs	Use case: Complex Data/Mobile	Use case: Internal Microservices

07. SECURITY: HTTPS & HASHING

The TLS Handshake Sequence



Salting and Hashing

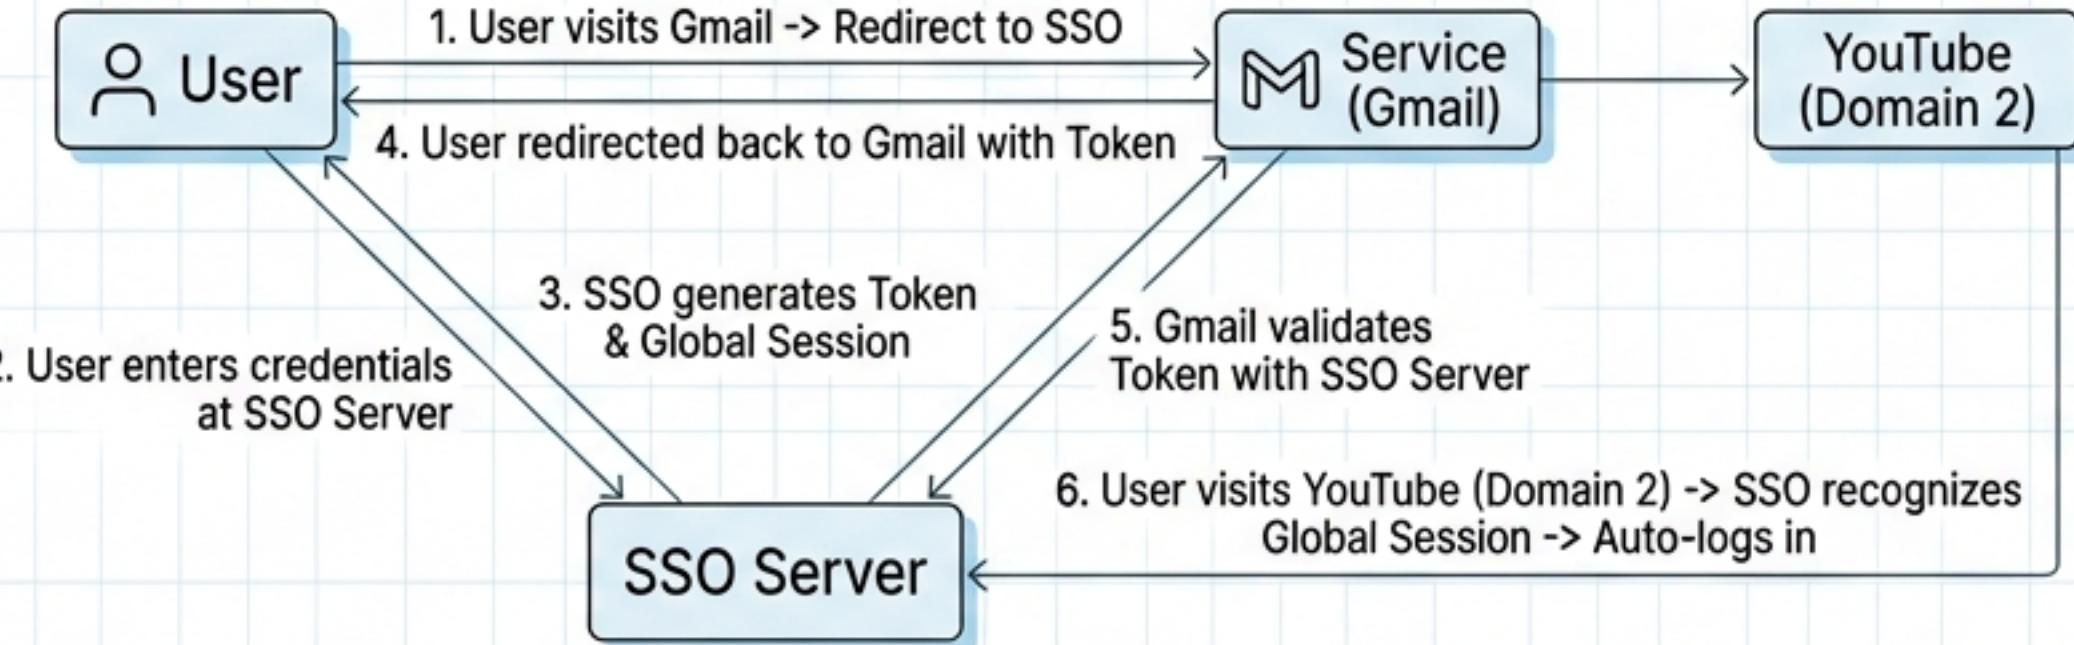


08. IDENTITY: SSO & HMAC

Helvetica Now Display

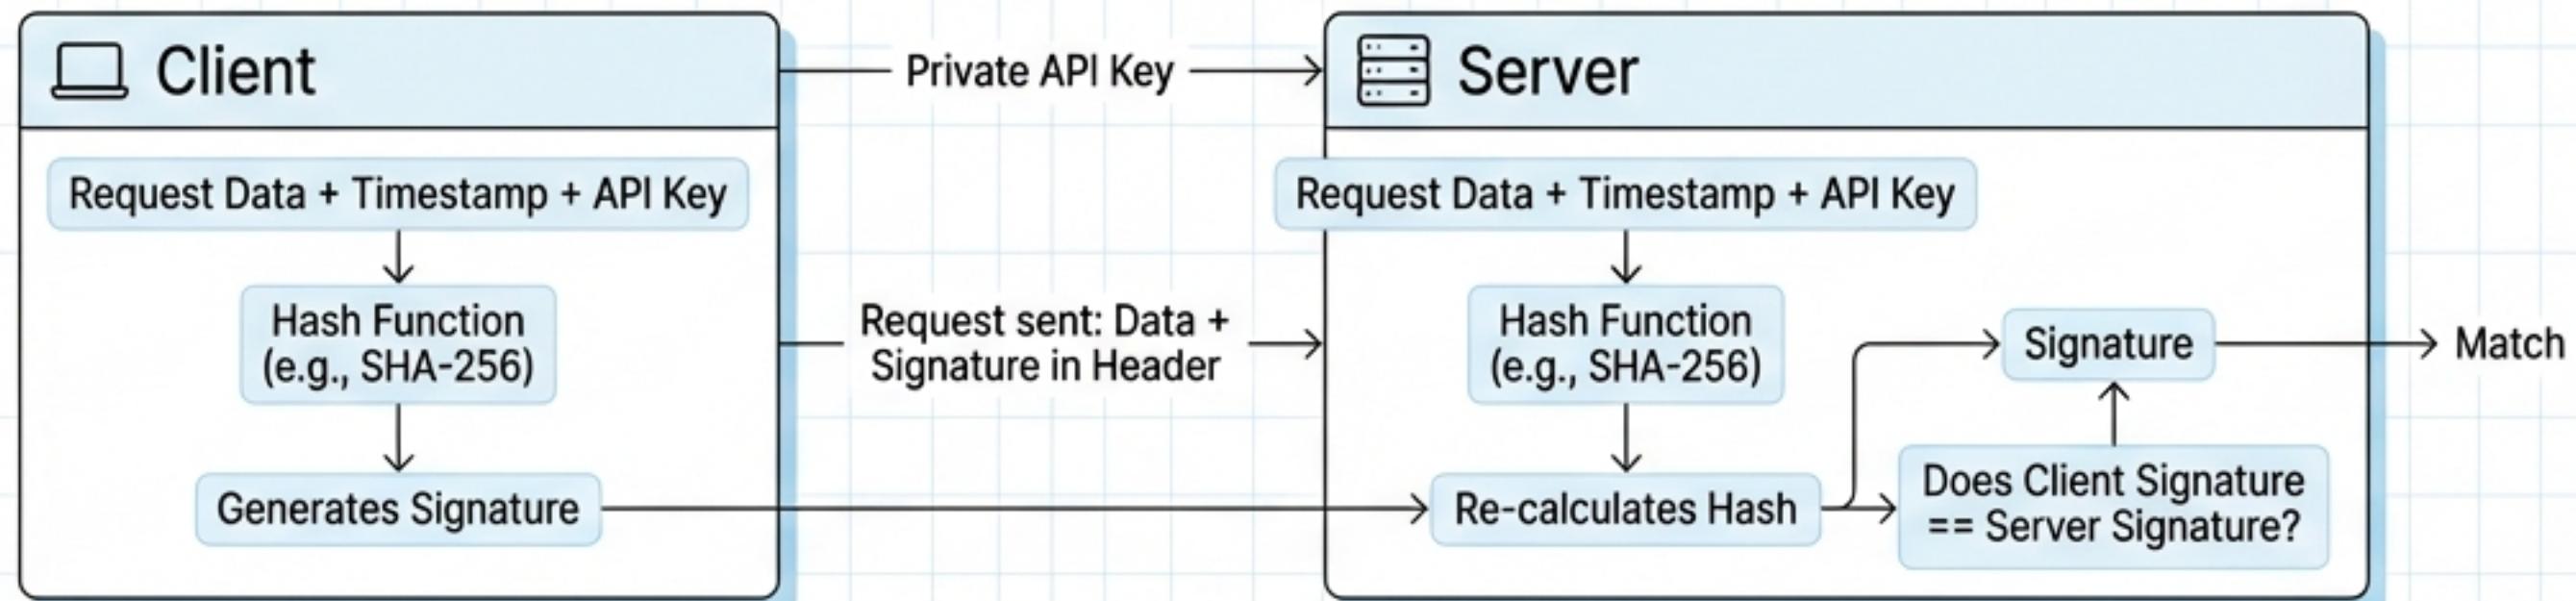
Single Sign-On (SSO) Flow

in Inter

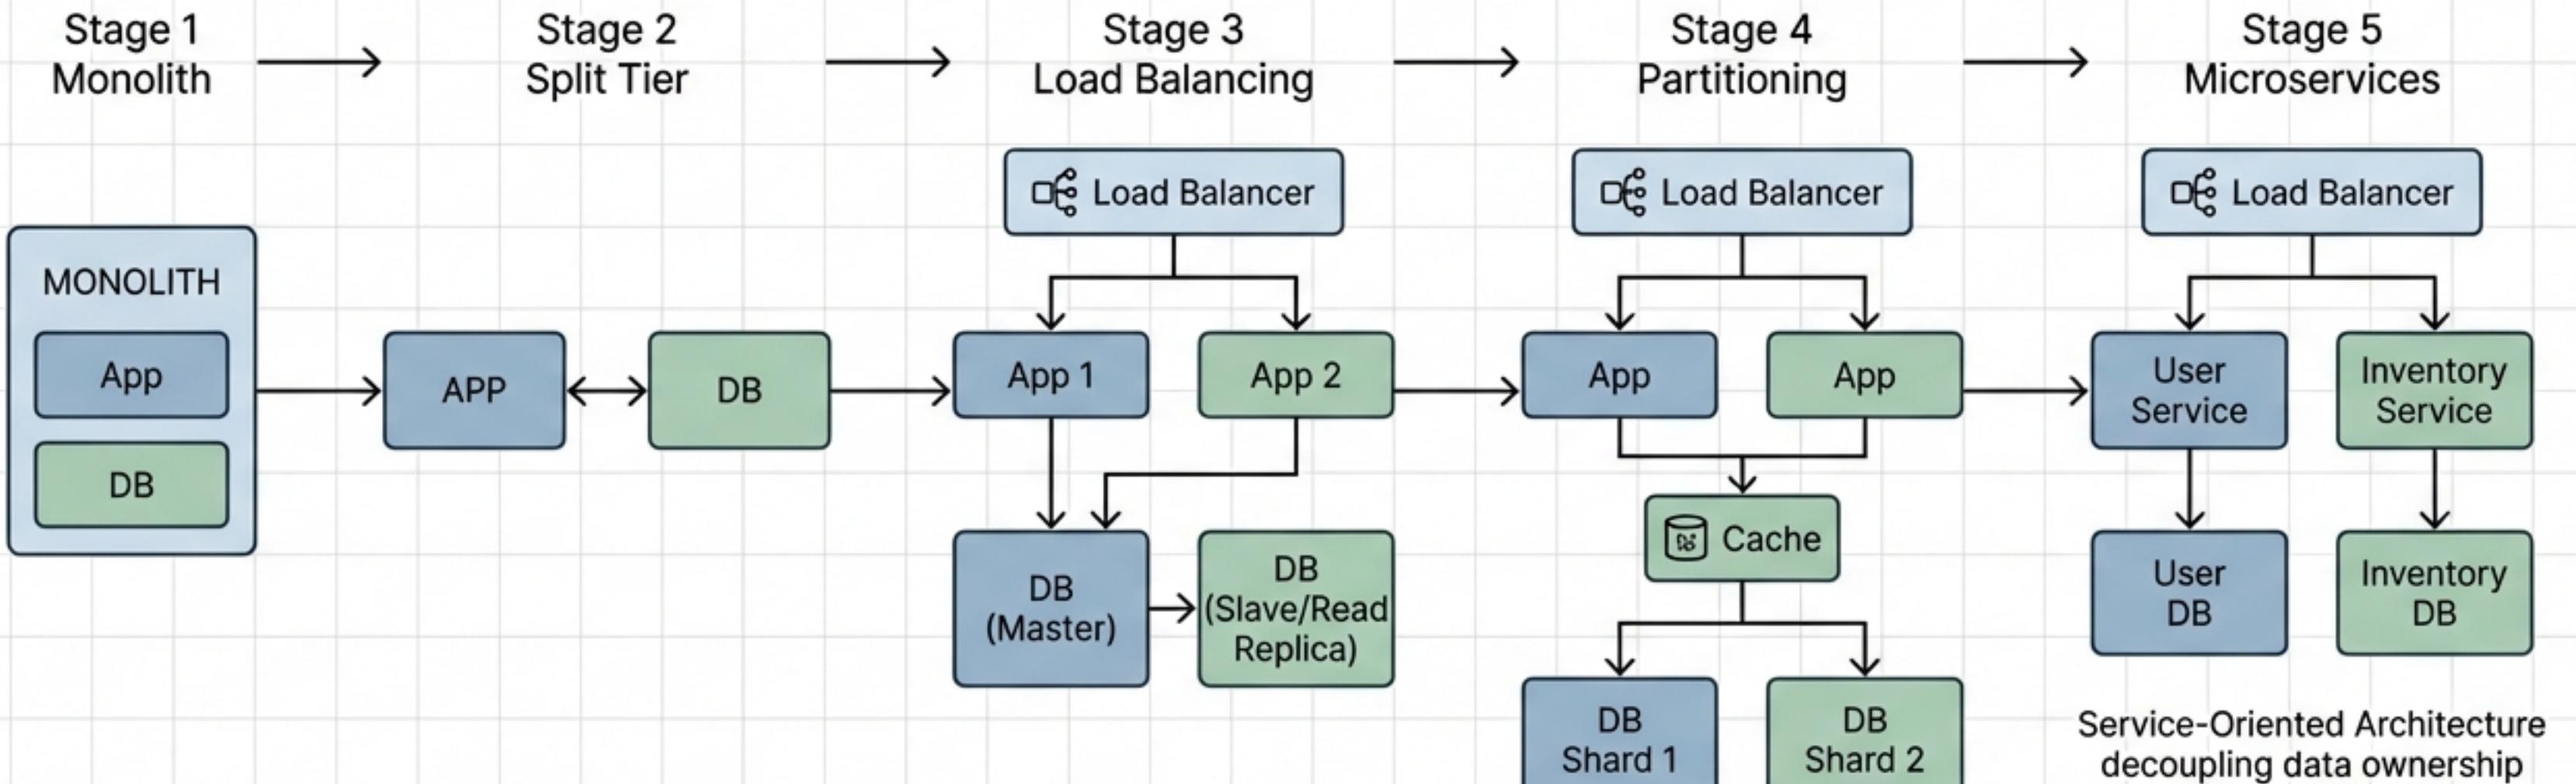


HMAC Authentication

in Inter



09. ARCHITECTURE: SCALING EVOLUTION



10. DATA PARTITIONING STRATEGIES

Helvetica Now Display

Vertical Partitioning

in Inter

User Table			
ID	Name	Bio	PhotoBlob
...
...
...
...
...

Vertical: Splitting by columns/feature.



User Core	
ID	Name
...	...
...	...

User Profiles		
ID	Bio	PhotoBlob
...
...

Horizontal Sharding

in Inter

User Table			
ID	Name	Bio	PhotoBlob
...
...
...
...
...
...

Rows 1 - 1,000,000

ID	Name	Bio	PhotoBlob
...
...
...
...
...
...

Horizontal Sharding

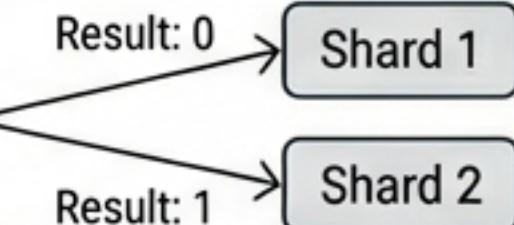
IDs 1 - 500,000			
ID	Name	Bio	PhotoBlob
...
...

IDs 500,001 - 1,000,000			
ID	Name	Bio	PhotoBlob
...
...

Shard 1
(Range Based)

Shard 2
IDs 500,001 - 1,000,000

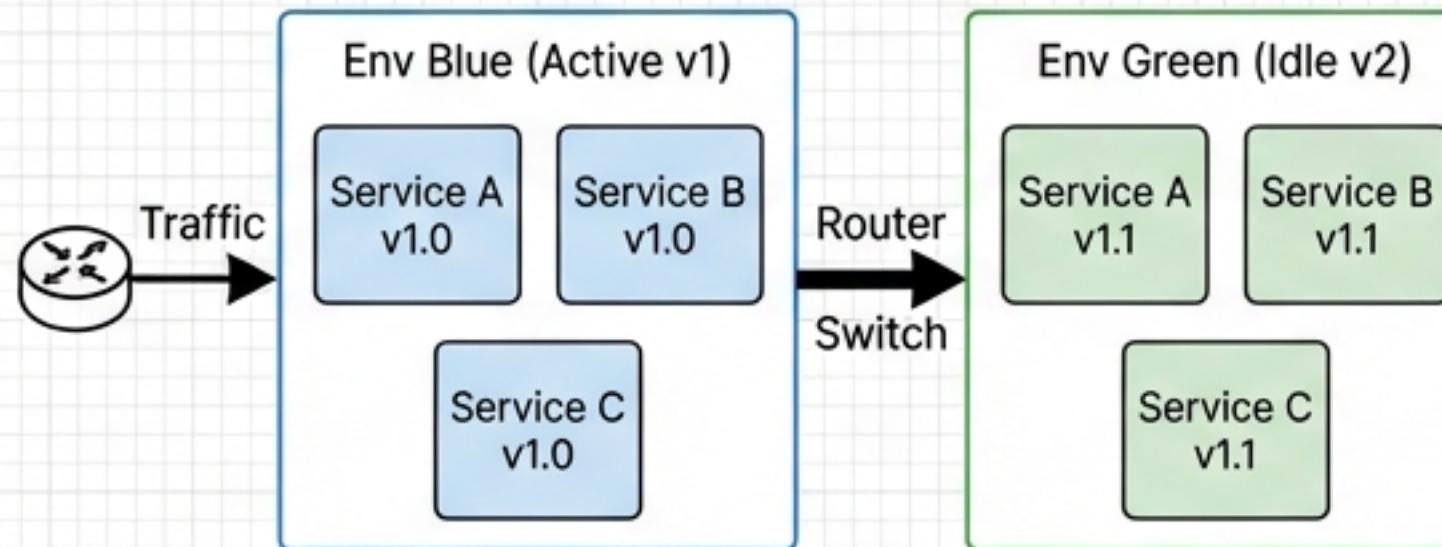
Hash Based Sharding
 $UserID \% 2$



11. DEPLOYMENT & COORDINATION

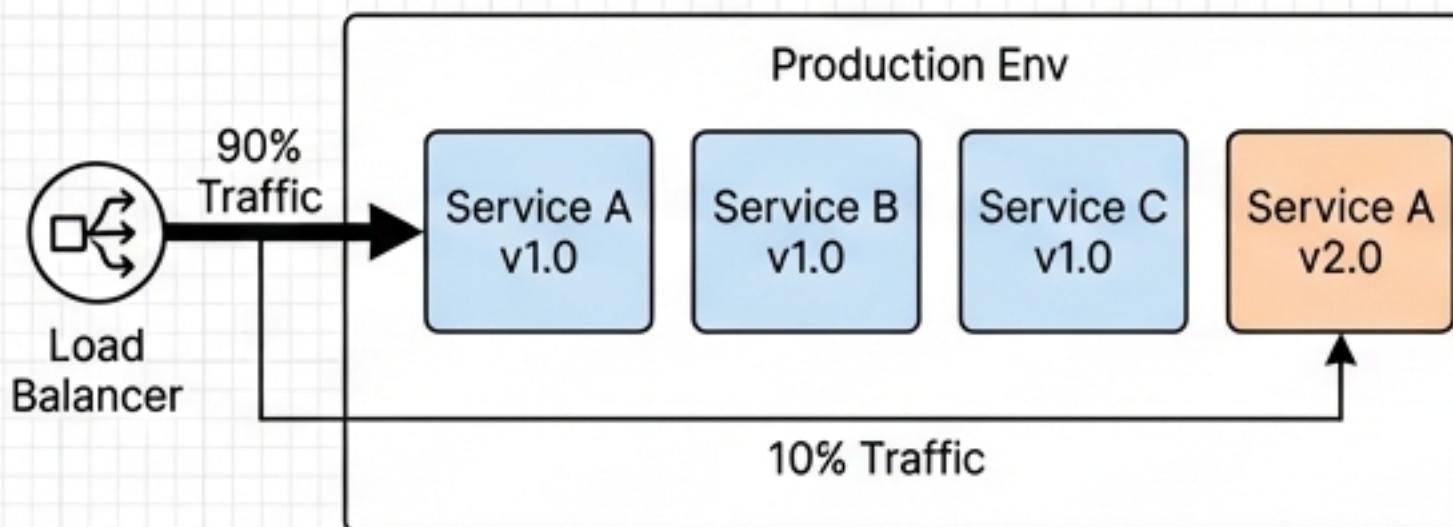
Deployment Strategies

Blue-Green Deployment



Zero-downtime deployment with instant rollback.

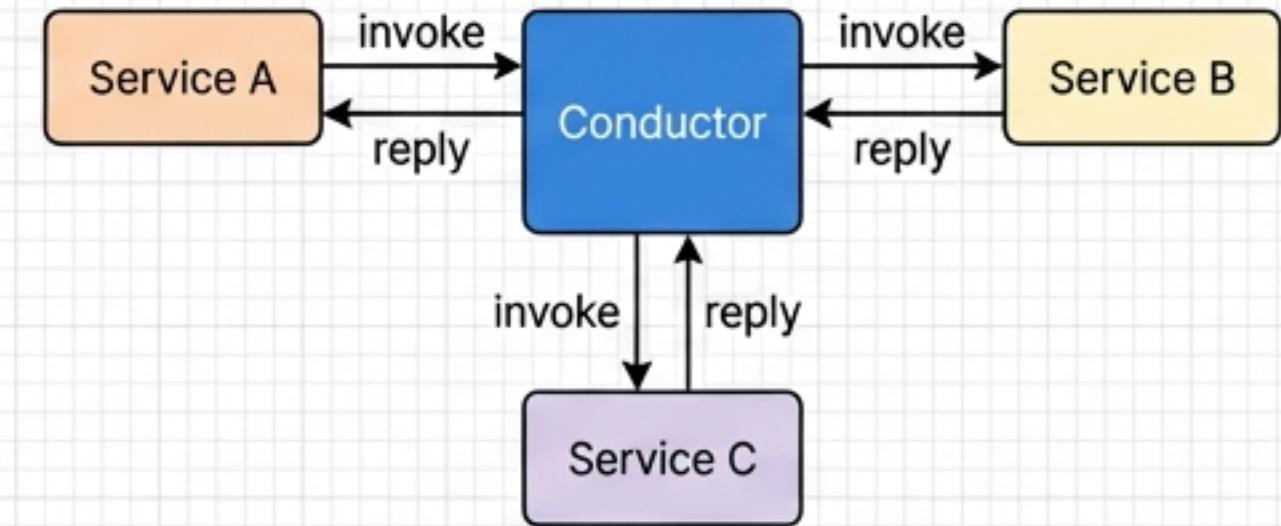
Canary Deployment



Rollout to a small user subset for testing.

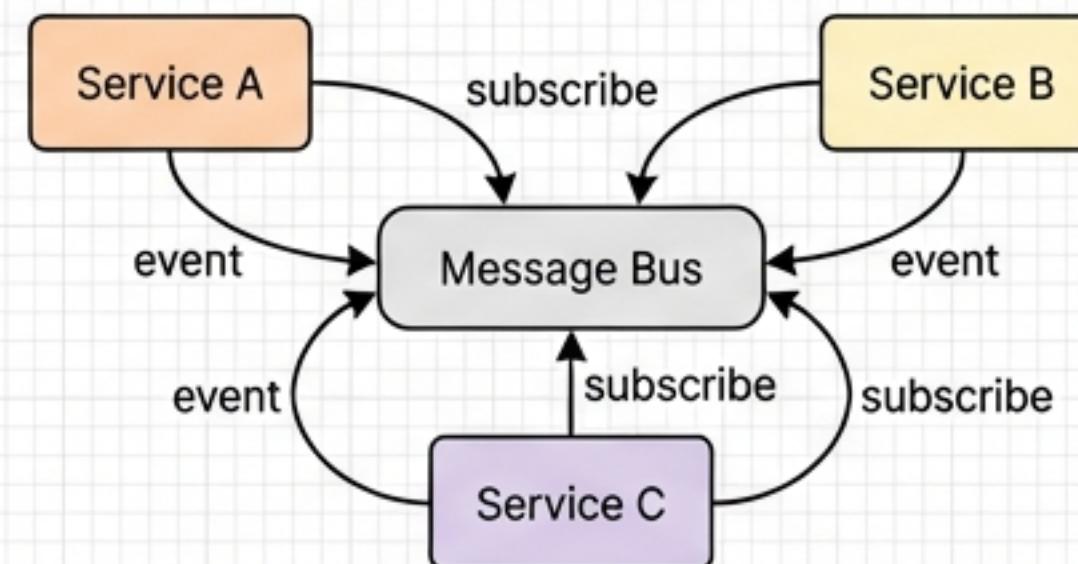
Microservices Interaction

Orchestration



Centralized control, tighter coupling.

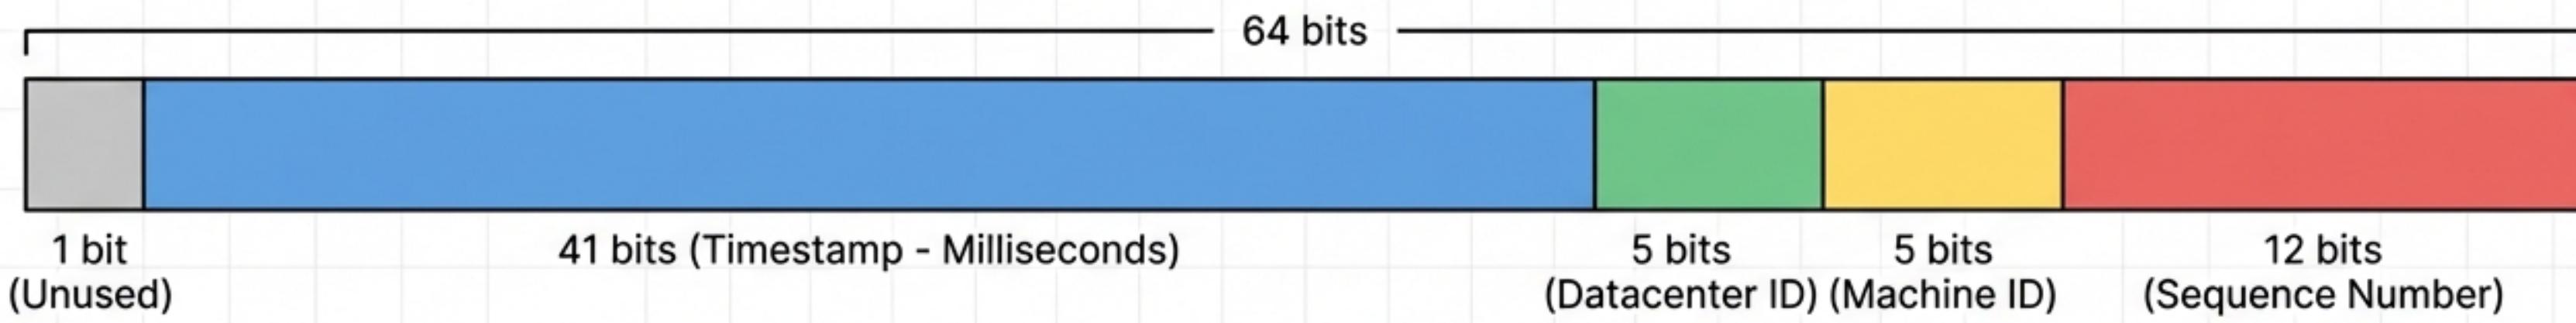
Choreography



Event-driven, loose coupling, asynchronous.

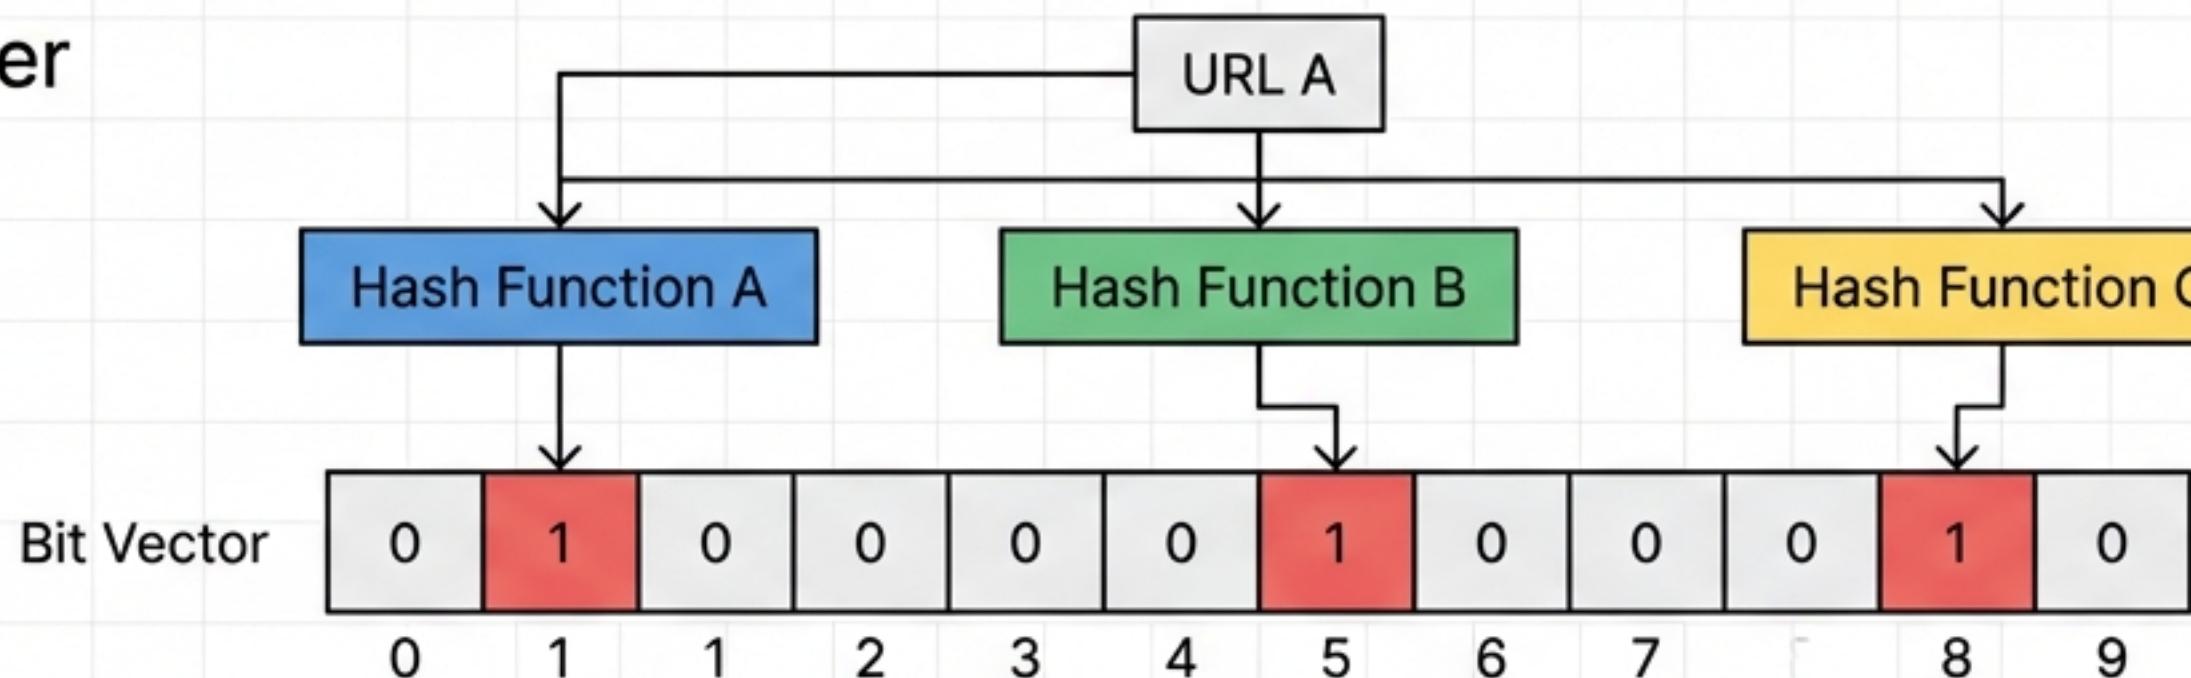
12. MECHANISMS: IDs & FILTERS

Snowflake ID Generator



Twitter Snowflake ID - Sortable & Distributed.

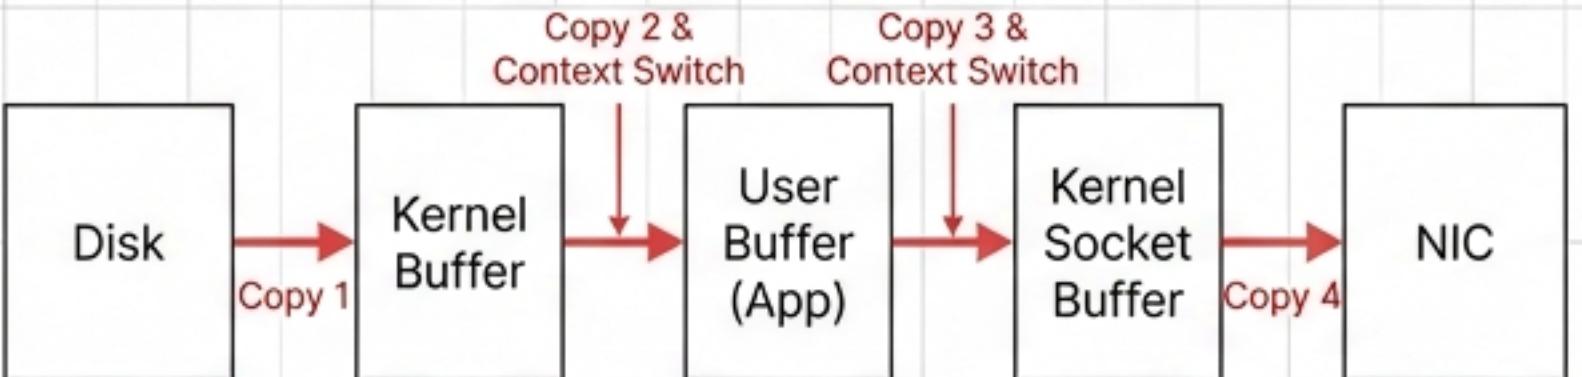
Bloom Filter



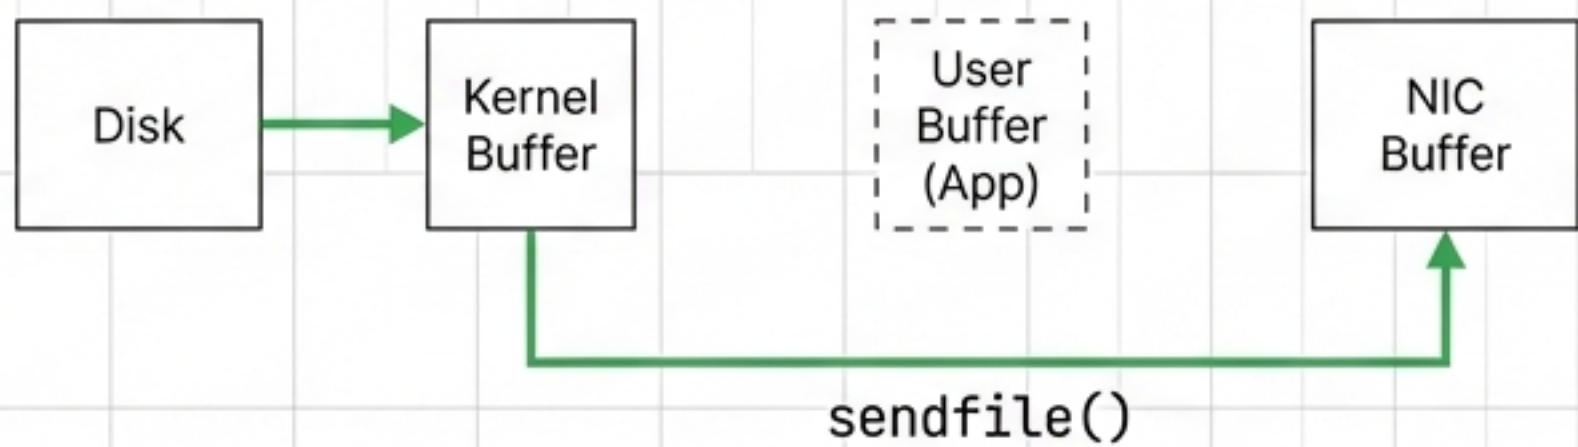
Probabilistic: If bits are 1, it MIGHT exist.
If any bit is 0, it DEFINITELY DOES NOT exist.

13. OPTIMIZATION: KAFKA ZERO COPY

Standard I/O (Inefficient)



Zero Copy (Efficient)



Requires 4 copies and 4 context switches.

Zero Copy reduces context switching by ~65%, allowing high throughput streaming.

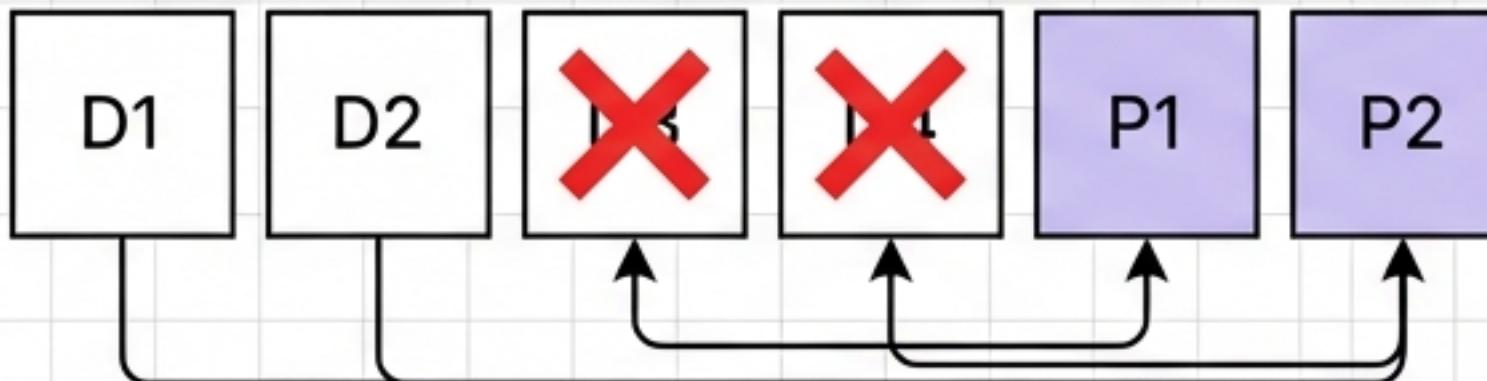
14. STORAGE: DURABILITY & UPLOADS

Erasure Coding



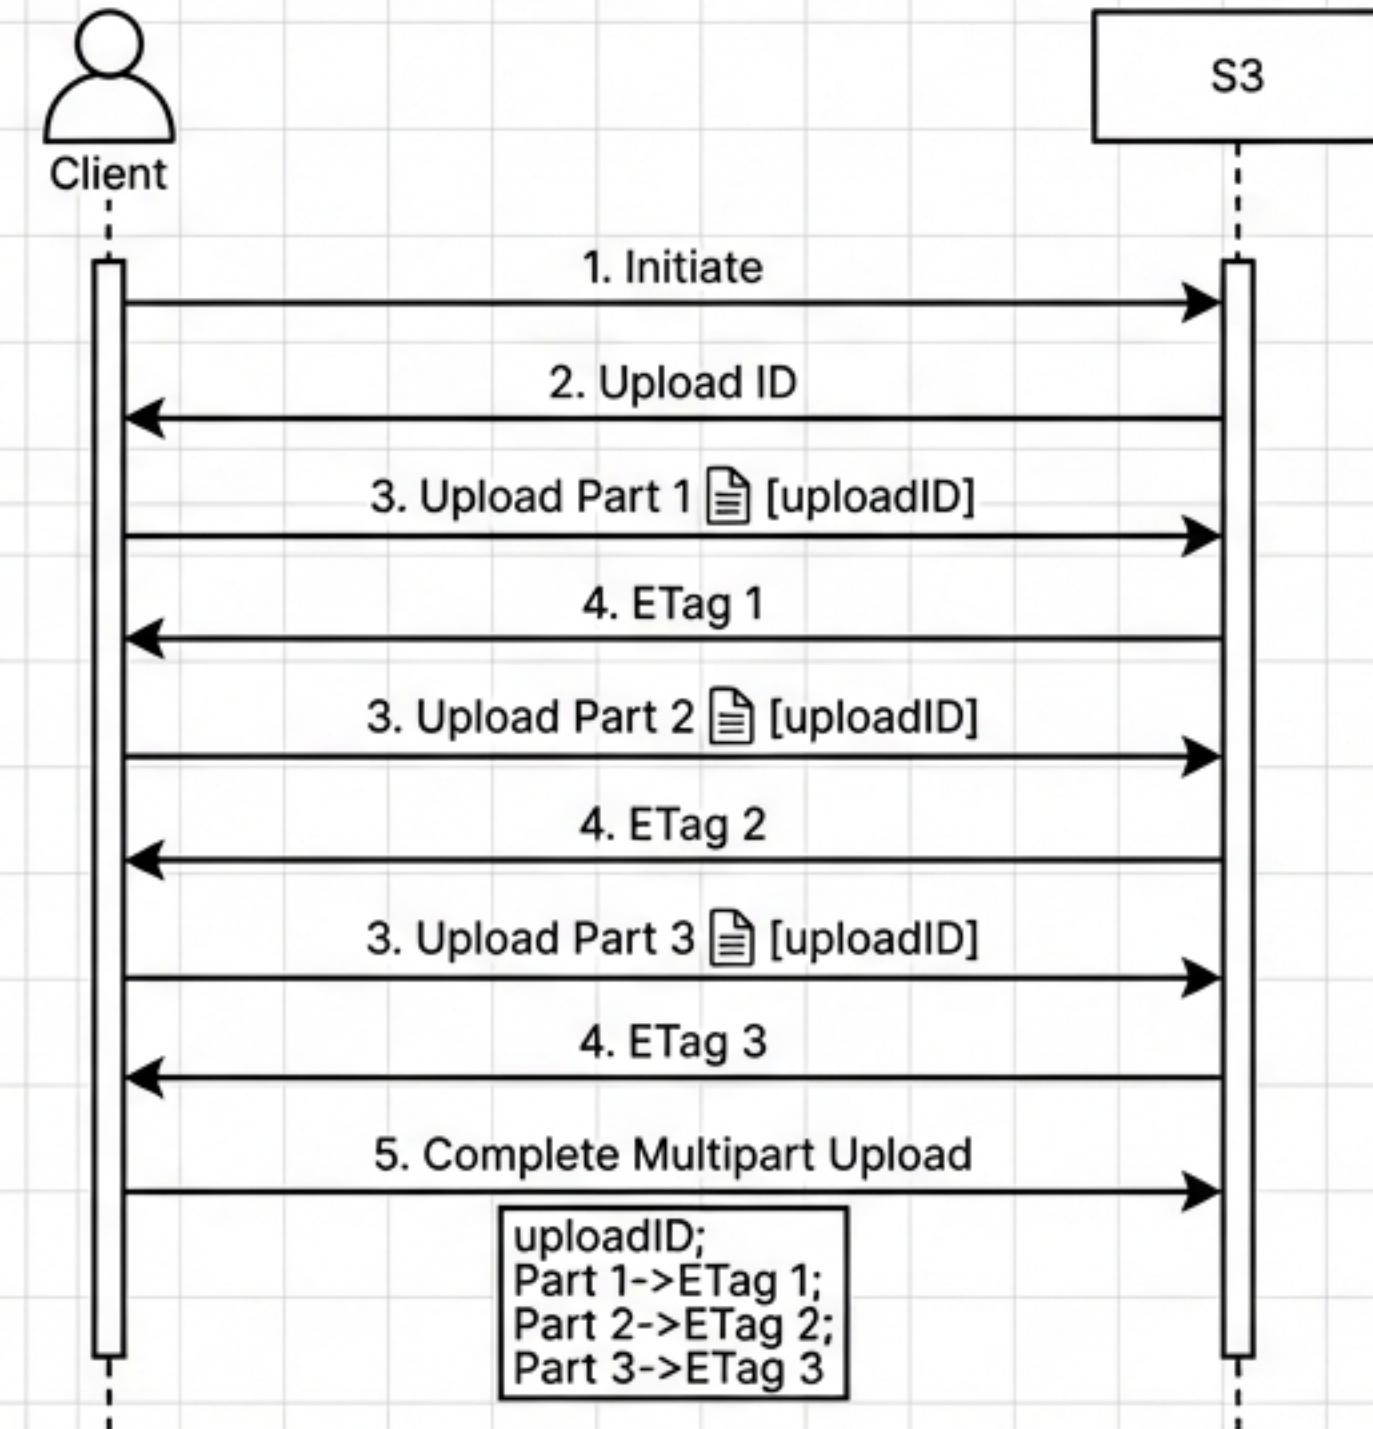
$$P1 = D1 + D2 + D3 + D4; P2 = f(D1, D2, D3, D4)$$

Data Reconstruction

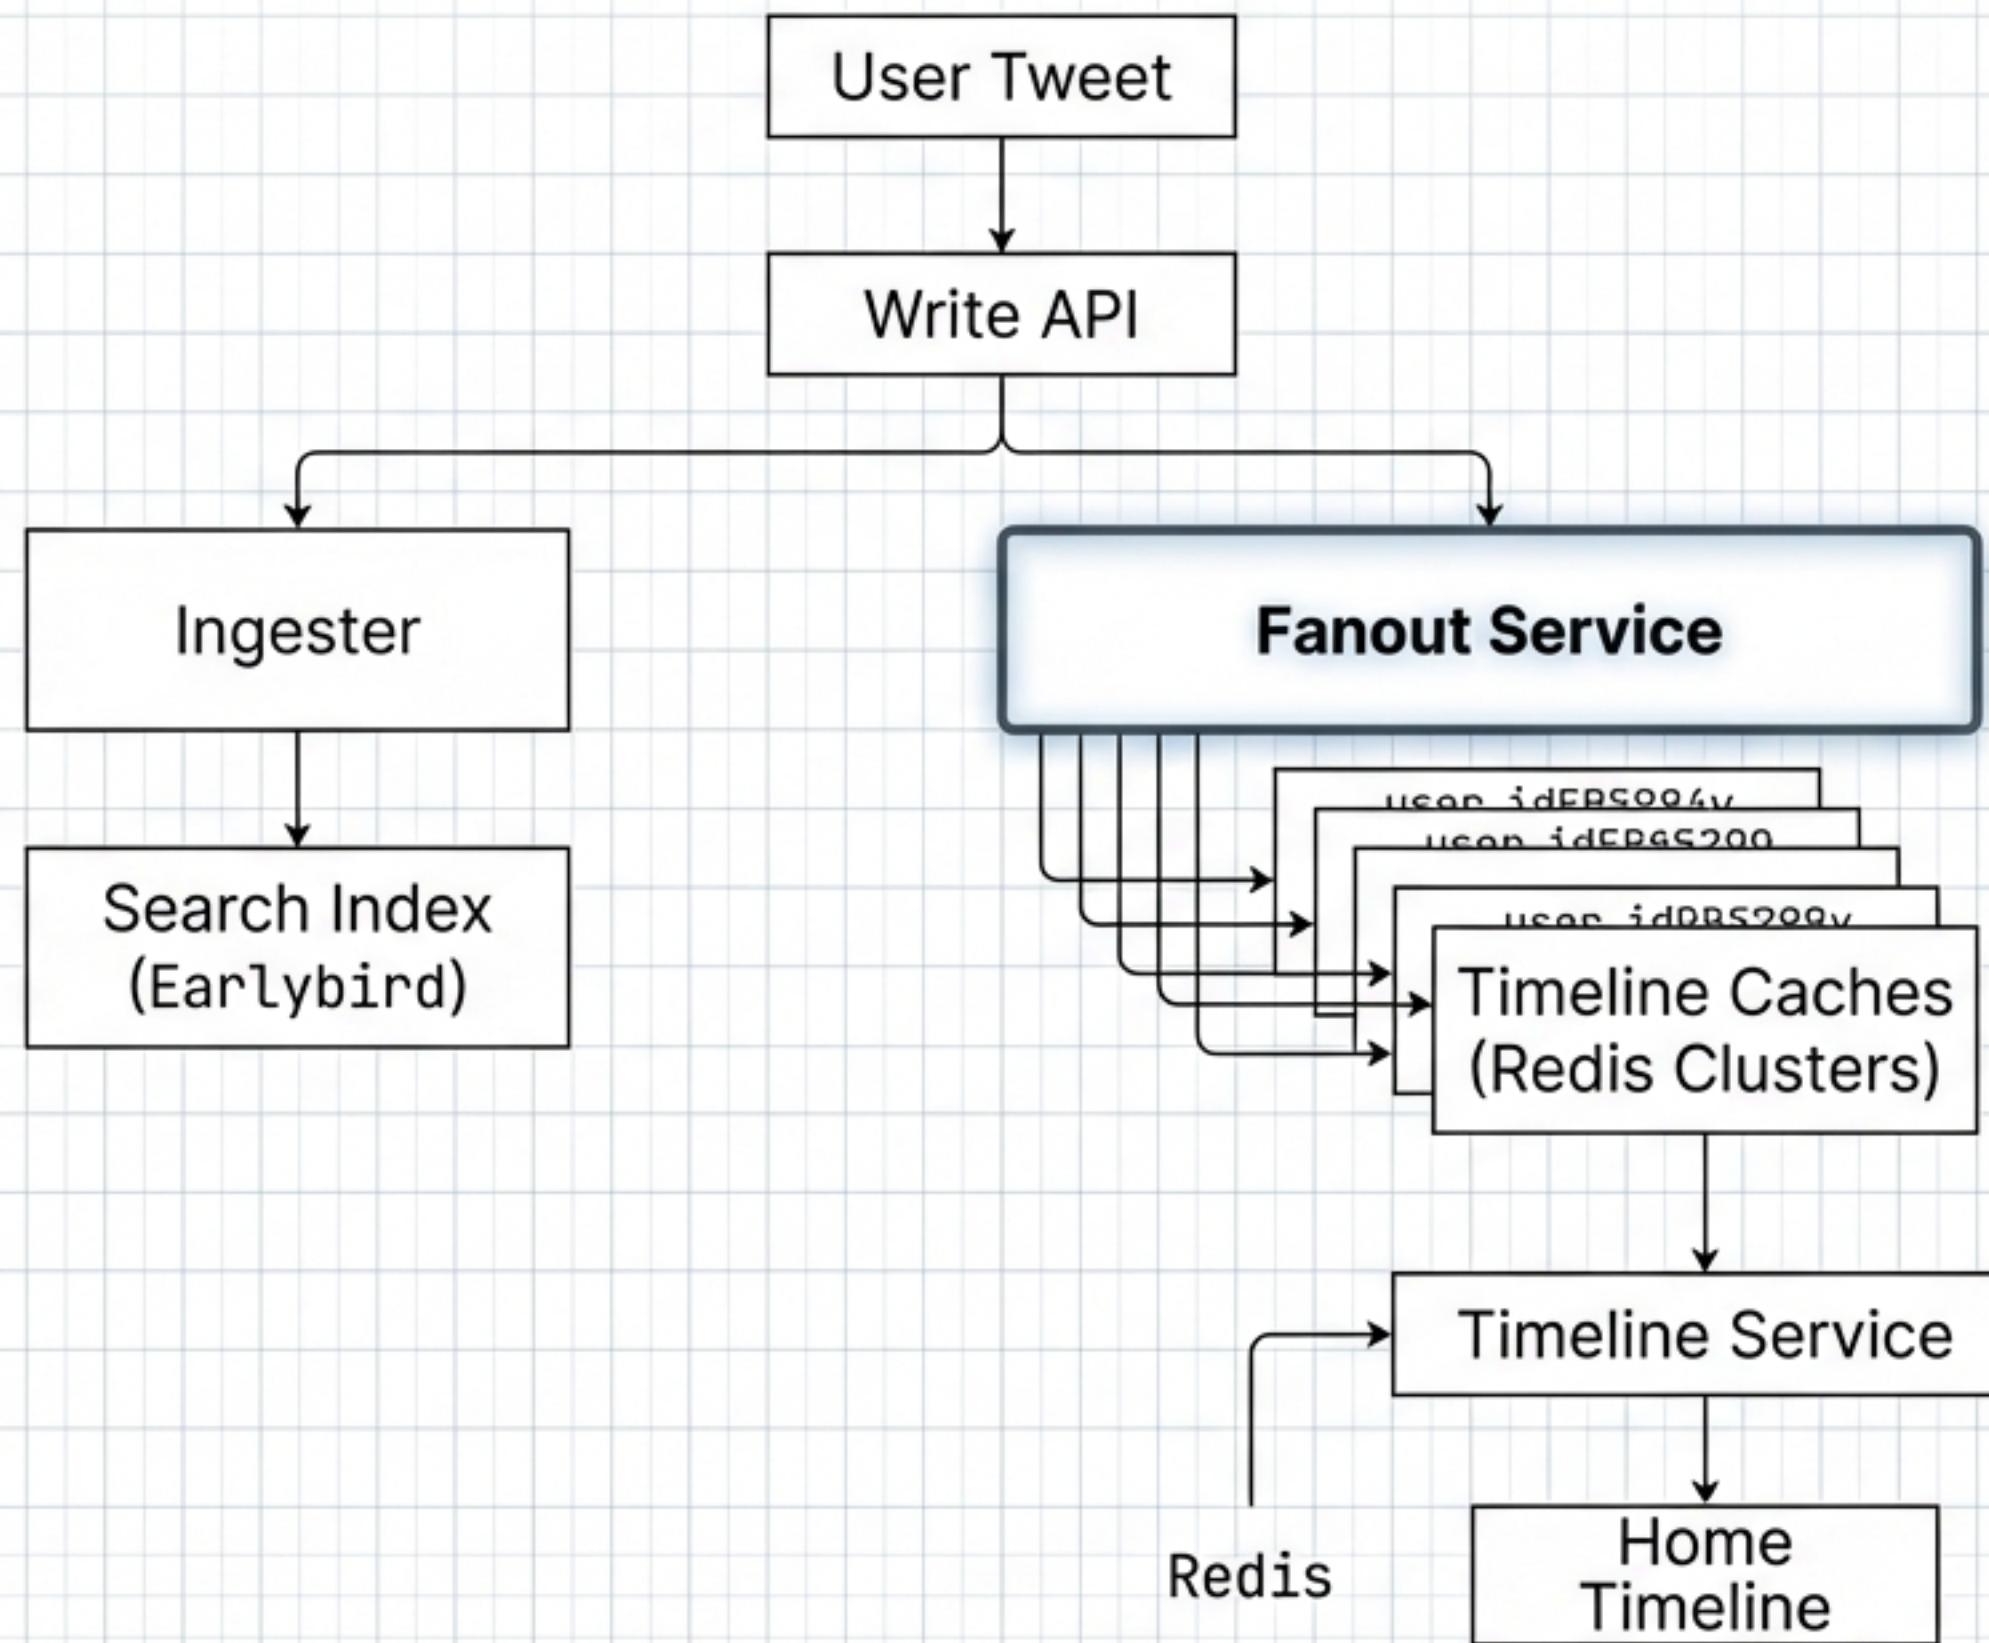


High Durability (11 nines) with only 50% storage overhead.

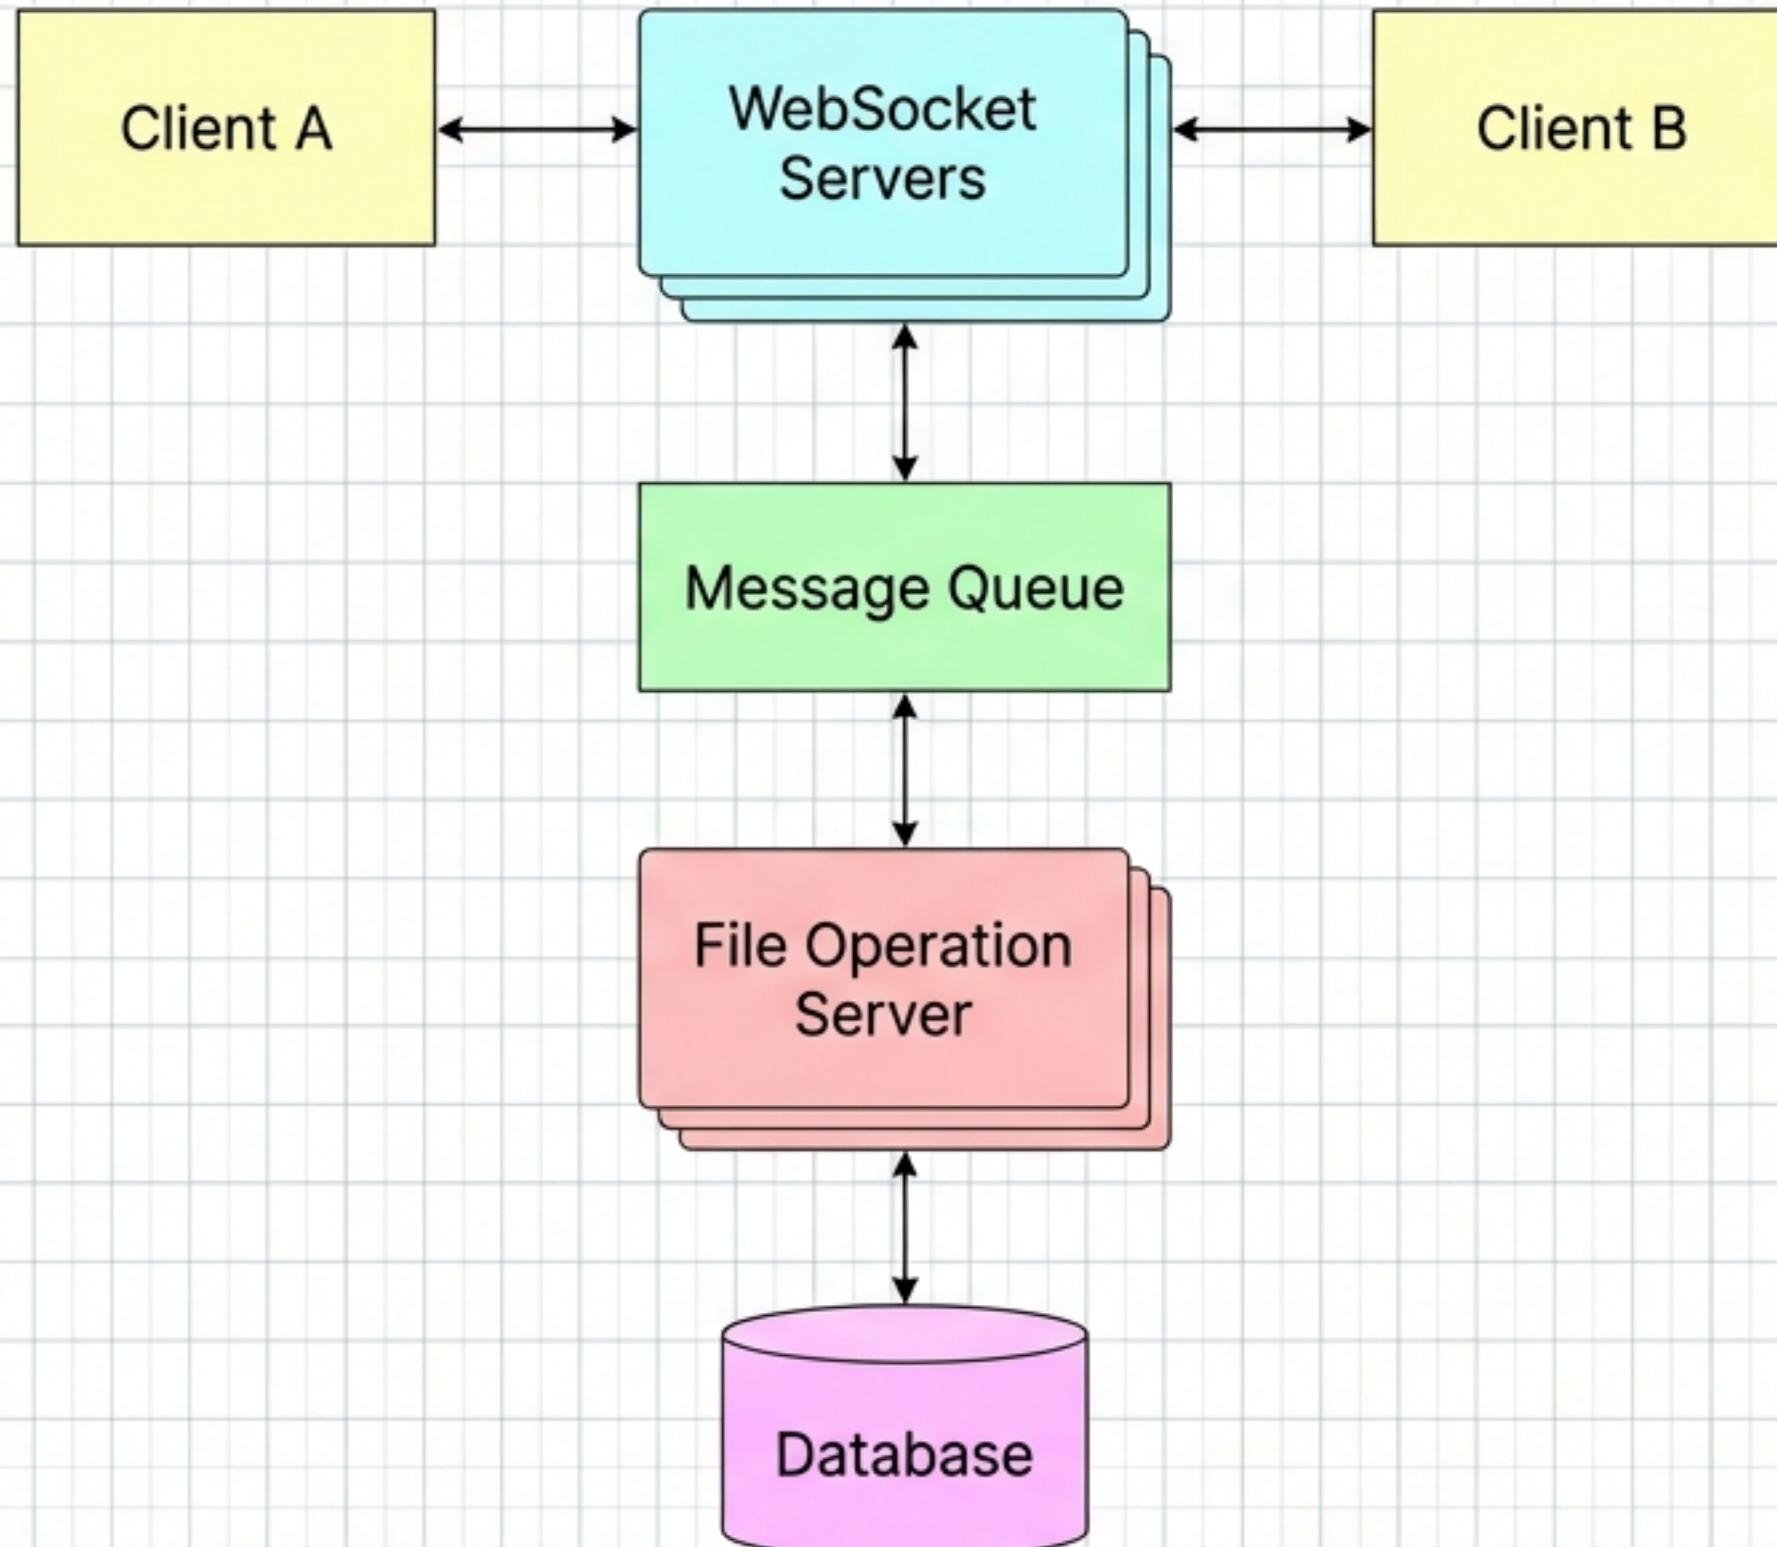
S3 Multipart Upload



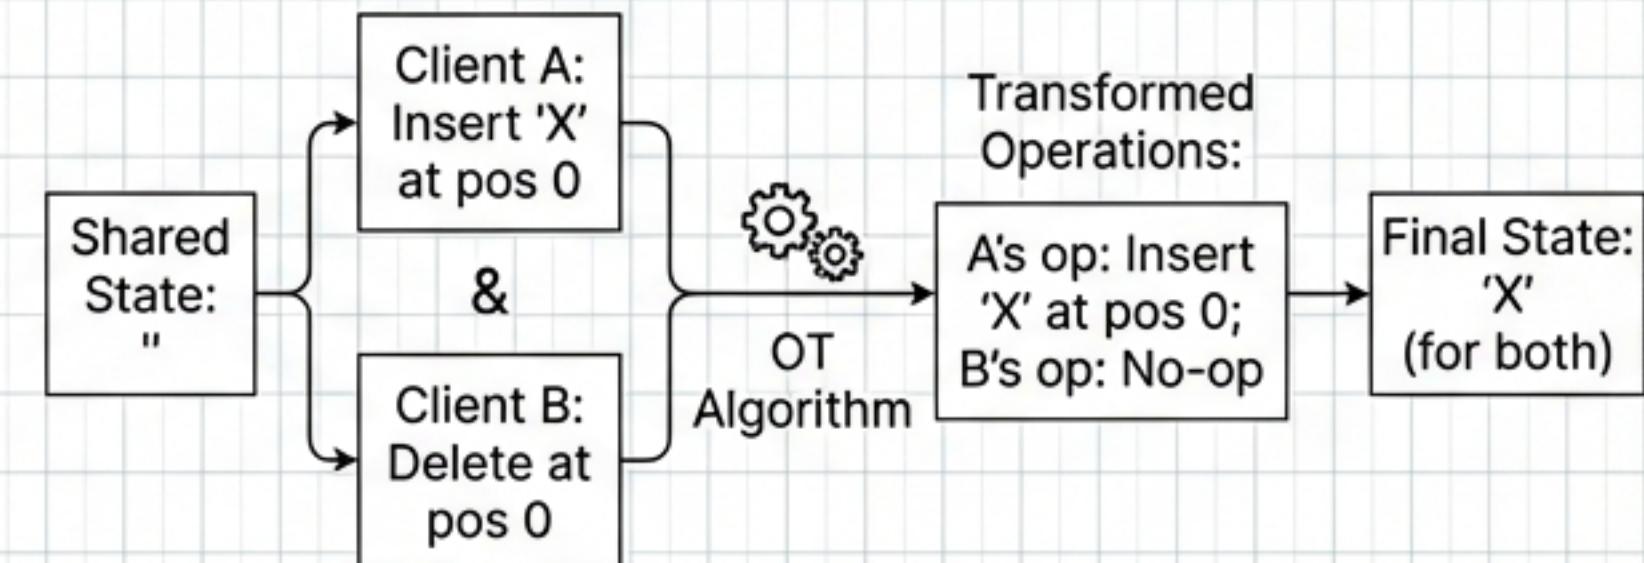
15. CASE STUDY: TWITTER ARCHITECTURE



16. CASE STUDY: GOOGLE DOCS (COLLABORATION)

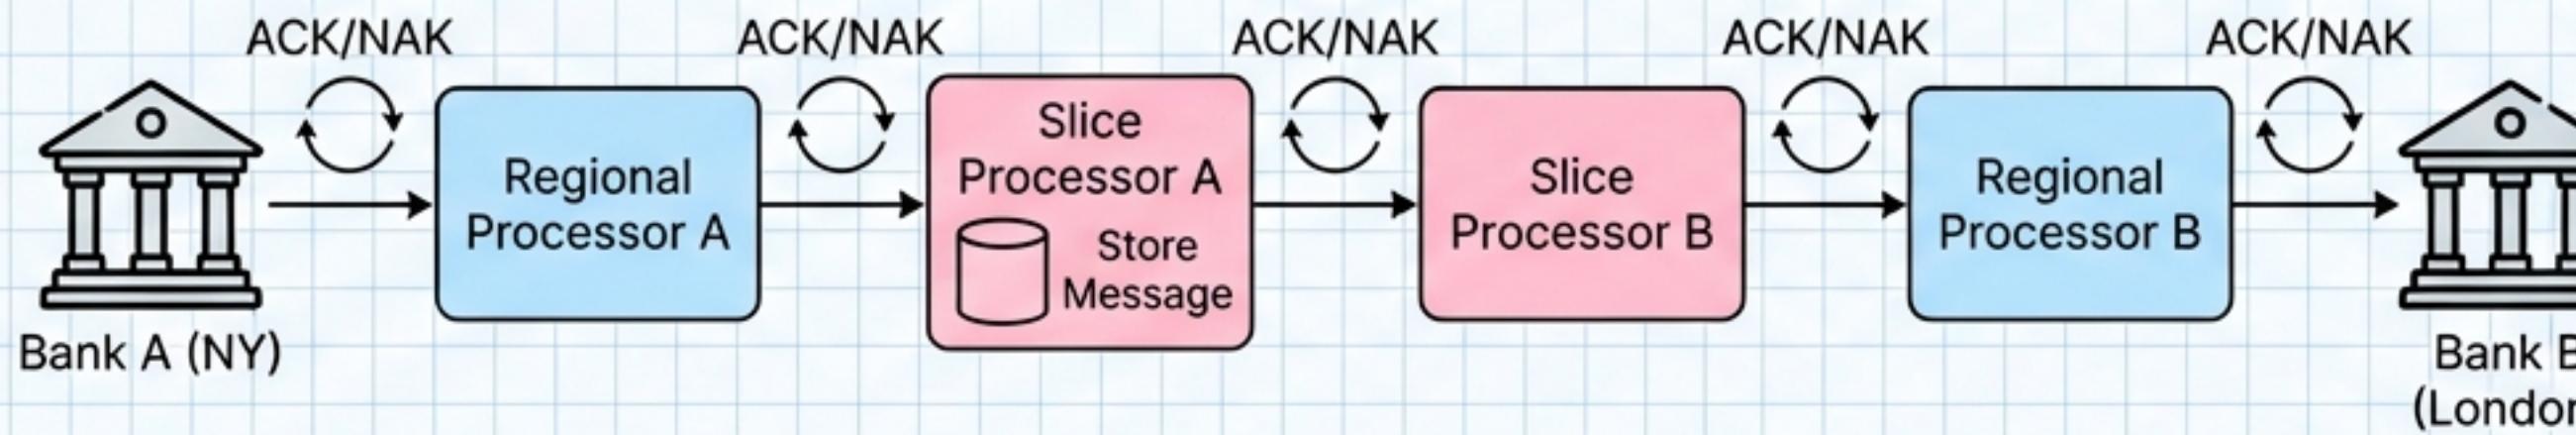


Operational Transformation (OT)



OT resolves concurrent editing conflicts
to ensure eventual consistency.
Code snippets's in JetBrains Mono.

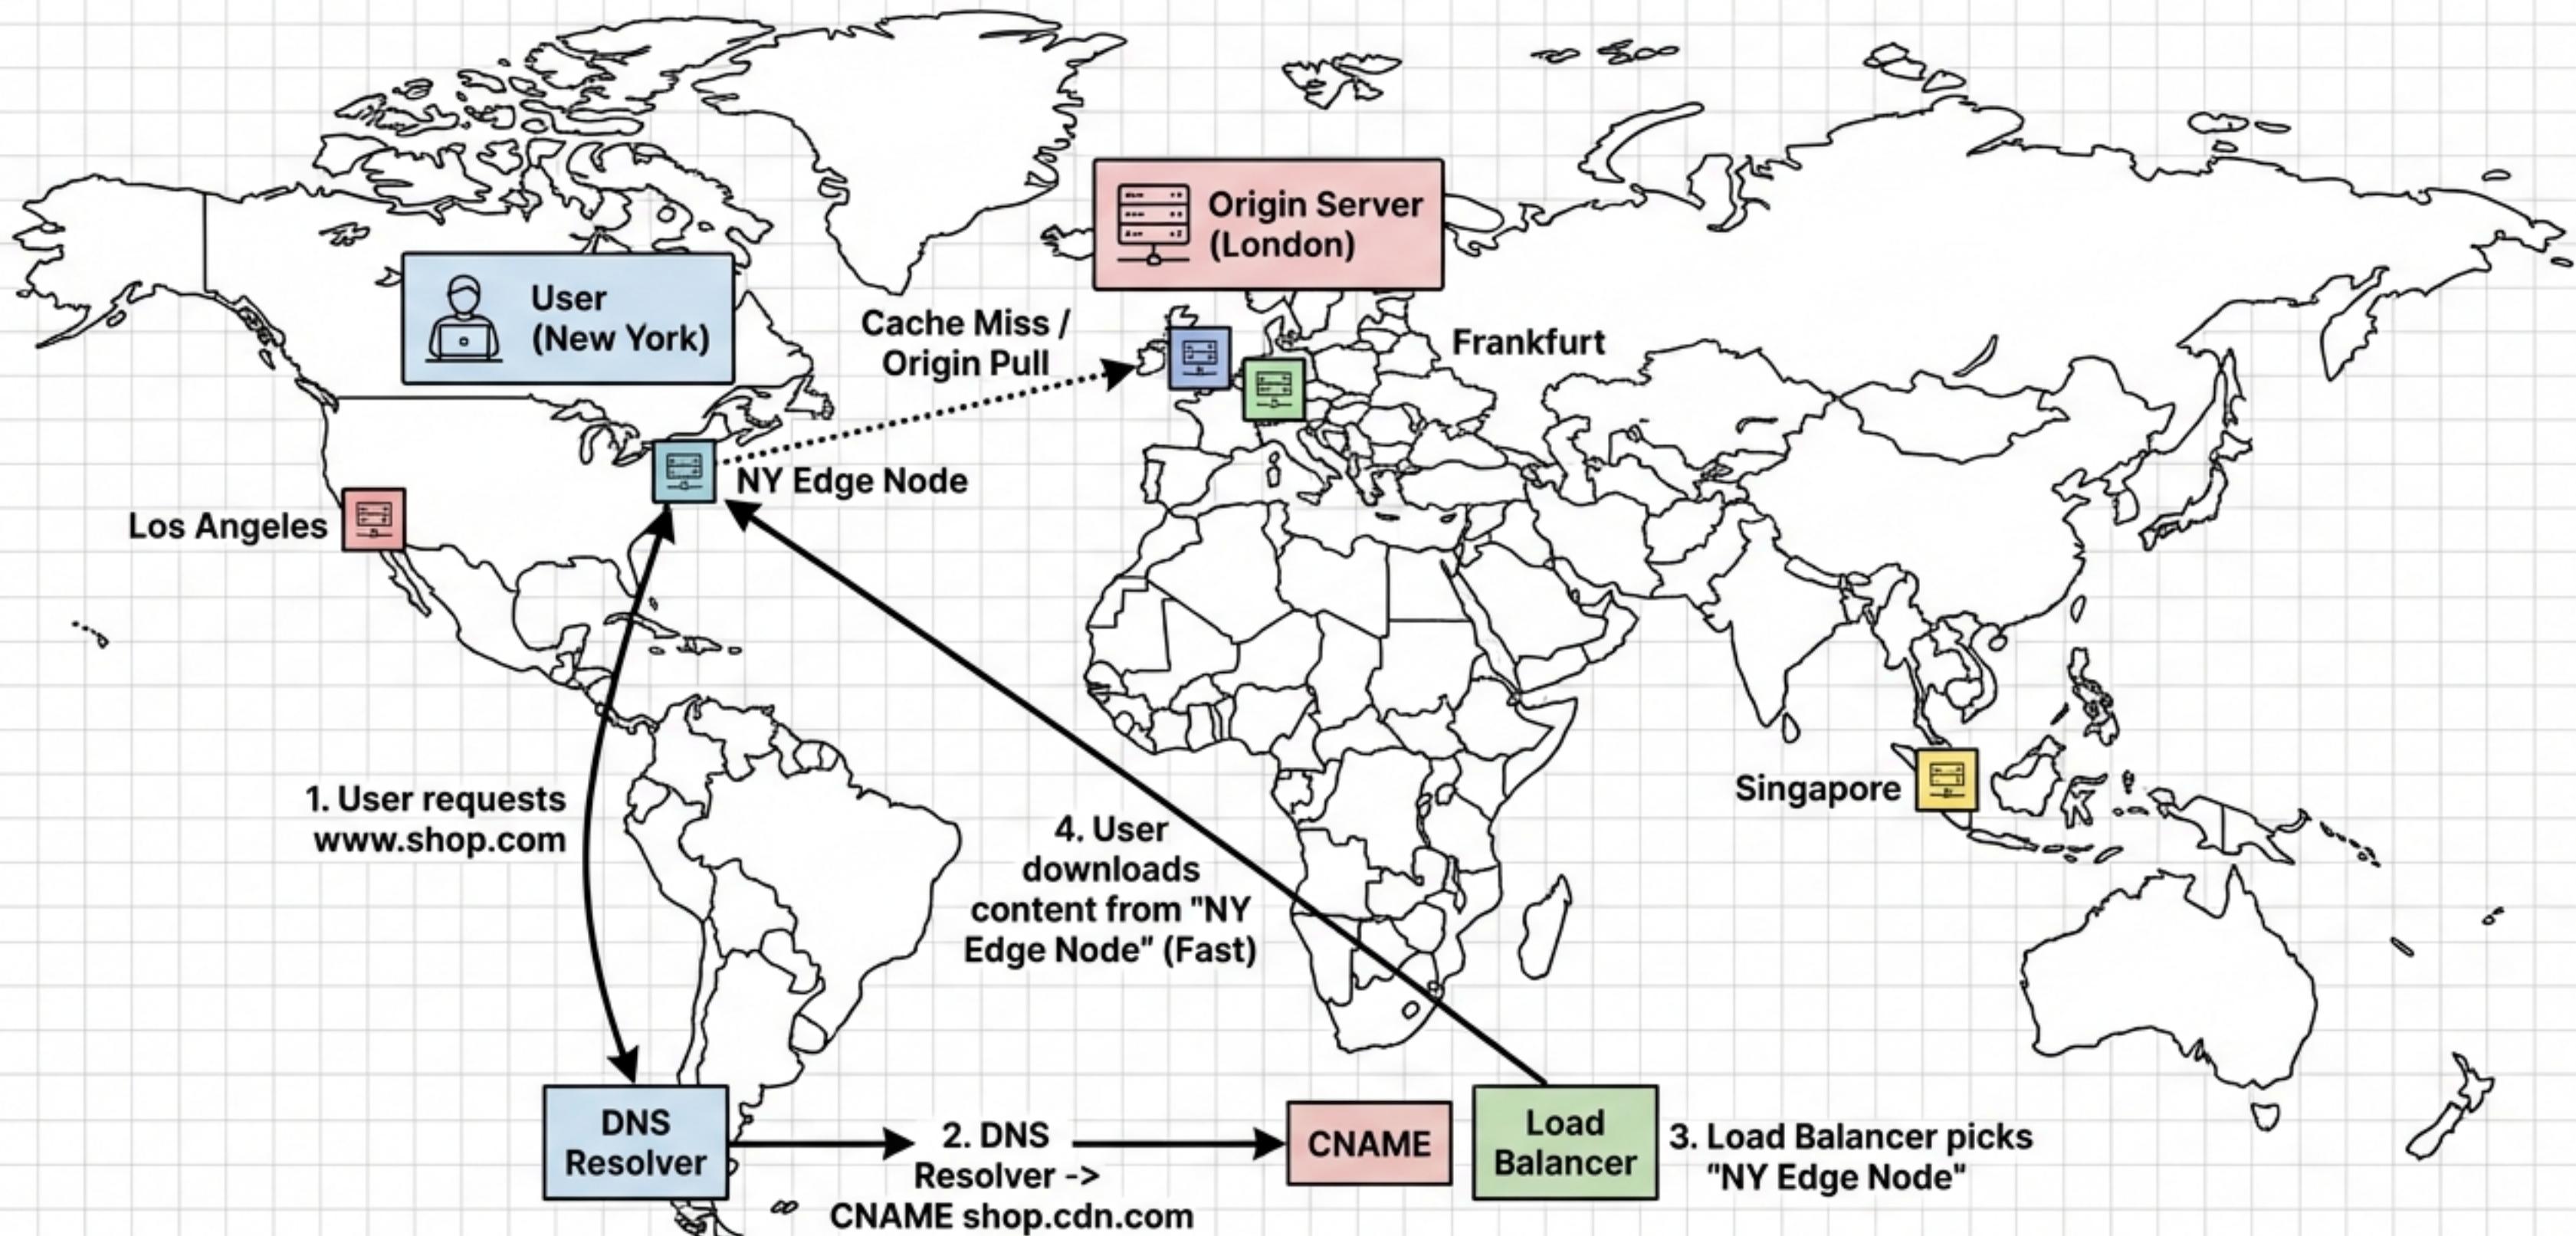
17. CASE STUDY: PAYMENTS & SWIFT



Key Patterns

- 1. Idempotency**
(Prevent double charge).
- 2. Reconciliation**
(Ledger matching).

18. GLOBAL SCALE: CDN & DNS



19. CONTINUOUS LEARNING



ARCHITECTURE

Inter Bold



- **The Amazon Builders' Library**

Inter Regular



- **Martin Fowler's Blog**

Inter Regular

DEVOPS & RELIABILITY

Inter Bold



- **The Phoenix Project**

Inter Regular



- **Google SRE Book**

Inter Regular

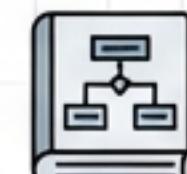


- **Accelerate**

Inter Regular

CODE & PATTERNS

Inter Bold



- **Head First Design Patterns**

Inter Regular



- **Designing Data-Intensive Applications**

Inter Regular

“System design is not about memorizing solutions, but understanding trade-offs.
Latency vs. Consistency. SQL vs. NoSQL. Complexity vs. Maintainability.”

Inter Medium